

Natural gas elasticities and optimal cost recovery under heterogeneity: Evidence from 300 million natural gas bills*

Maximilian Auffhammer and Edward Rubin

November 6, 2017

Abstract

In 2016 natural gas became the United States' primary source of energy for electricity generation. In addition, more than 50% of American households heat primarily with natural gas. Hence understanding residential natural gas consumption behavior has become a first-order problem. In this paper, we provide the first ever causally identified, microdata-based estimates of residential natural gas demand elasticities using a panel of approximately 300 million bills in California. To overcome multiple sources of endogeneity, we employ a two-pronged empirical strategy: (1) we exploit a discontinuity along the border between two major natural-gas utilities in conjunction with (2) an instrumental variables strategy based upon the differences in the utilities' rules/behaviors for internalizing changes in the upstream natural gas spot market. We estimate the elasticity of demand for residential natural gas is between -0.23 and -0.17 . We also provide evidence of seasonal and income-based heterogeneity in this elasticity. This heterogeneity provides unexplored policy avenues that may be simultaneously efficiency enhancing and pro poor.

*Corresponding author: Edward Rubin, 207 Giannini Hall #3310, University of California, Berkeley, CA 94720-3310. Email: edward@berkeley.edu. Auffhammer: auffhammer@berkeley.edu. We thank the Pacific Gas and Electric Company, the Southern California Gas Company, and San Diego Gas & Electric for making their billing data available. We thank Reed Walker, Sofia Berto Villa-Boas, James Sallee, Sol Hsiang, Meredith Fowlie, Lucas Davis, Severin Borenstein, Peter Berck, and seminar/workshop participants at Berkeley ARE, Cornell University, the Lulea Department of Economics, Yale University, Camp Resources, and the Heartland Environmental and Resource Economics Workshop for valuable feedback. All remaining errors are the authors.

Note: Please do not cite without permission from the authors.

1 Introduction

Throughout the late 19th and most of the 20th century, coal dominated all other fossil fuels, powering the unprecedented economic transformation of the United States and many other major economies. However, due to the invention and broad rollout of a new technology to extract gas from below the surface—hydraulic fracturing (“frac(k)ing”)—it appears as though ample supplies of low cost natural gas will power the first half of the 21st century. Additionally, in 2005, hydraulic fracturing received significant exemptions from the Clean Air Act, the Clean Water Act, and the Safe Drinking Water Act via the Energy Policy Act of 2005 [Environmental Protection Agency 2013]. Since 2005, natural gas production in the United States has expanded dramatically, and natural gas prices have fallen considerably—often residing at half of their pre-2005 levels (see Figure 3 and Hausman and Kellogg 2015). In 2016, natural gas surpassed coal as the main source of energy for electricity generation in the United States U.S. Energy Information Administration 2016b. Further, natural gas is the main heating fuel for more than half of US residences. The US Energy Information Administration (EIA) estimates that energy-related carbon dioxide (CO₂) emissions from natural gas overtook energy-related CO₂ emissions from coal in 2016 [U.S. Energy Information Administration 2016c].

The low price of natural gas and its recently abundant volumes, coupled with natural gas’s status as the cleanest and most efficient fossil fuel [National Academy of Sciences 2016; Levine, Carpenter, and Thapa 2014], have certainly garnered broad public and policy support for natural gas.¹ Due to its low carbon content per BTU, natural gas is often regarded as a “bridge fuel”—bridging society toward a future mainly powered by largely carbon-free sources of renewable energy. However, natural gas is not without critics—particularly natural gas that results from hydraulic fracturing. The most common criticisms of current natural gas policy relate to environmental degradation—including groundwater contamination, triggering small earthquakes, increases in air pollution, and accidents from the large number of trucks servicing fracking sites. More broadly, researchers have also critiqued inefficient and potentially regressive pricing (and regulatory) regimes used in the consumer-facing side of the industry

¹The fact that an increasingly large share of natural gas is produced in the United States also wins natural gas considerable political support [Levine, Carpenter, and Thapa 2014].

[Borenstein and Davis 2012; Davis and Muehlegger 2010].

In this paper, we provide the first causally identified, micro-data based estimates of the elasticity of demand for residential natural gas. This parameter estimate contributes to a body of natural gas research requiring a precise and well-identified estimate of the price elasticity of demand of natural gas; these topics range from the welfare implications of the United States' natural gas boom to hydraulic fracturing's environmental impact to the efficiency and distribution of natural gas's prevailing pricing and regulatory standards [Davis and Muehlegger 2010; Hausman and Kellogg 2015]. Despite this policy relevance, there is a relative dearth of well-identified estimates for the own-price elasticity of the demand for natural gas. A cursory Google Scholar search returns approximately 148,000 results related to *economics, elasticities, and electricity*; equivalent searches for *coal* and *gasoline* return approximately 70,000 results each. A similar search for articles related to *natural gas* finds fewer than 40,000 results.² Perhaps more importantly, we are unable to find any published research that pairs consumer-level data with appropriate identification strategies to estimate a price elasticity of demand for natural gas that carries a causal interpretation.

This paper draws upon a dataset of approximately 300 million residential natural gas bills in California, which provides us with the requisite power and detail to resolve multiple sources of endogeneity and identify both pooled and heterogeneous price elasticities. In order to avoid the potential biases from the sources of endogeneity discussed below, we employ a two-pronged identification strategy. The first prong of our identification strategy exploits a spatial discontinuity based upon the boundary between the service areas of two large natural gas utilities. The second prong interacts this spatial discontinuity with a *supply-shifting* instrumental variables (IV) approach—instrumenting the utilities' consumer-facing prices with the weekly average spot price of natural gas at a major natural gas distribution hub in Louisiana (the *Henry Hub*). As a result of this two-part empirical strategy—coupled with the rich set of fixed effects that our dataset allows—the identifying variation in the residential price of natural gas comes from (1) on which side of a long-established border between utility-owned natural gas networks the household is located, and (2) the subtly different pricing rules governing the two rate-of-return-earning utilities as they individually

²The authors performed these searches in January 2017.

respond to price variation in the natural gas spot market. This approach yields an estimate for the price elasticity of demand for residential natural gas ranging from -0.23 to -0.17 .

In addition, we find evidence of heterogeneity in this elasticity along two economically important and statistically significant dimensions: season and income. Lower-income households and higher-income households both are essentially inelastic to price in summer months. In winter months, however, lower-income households are substantially more elastic to prices than higher-income households. In addition to providing motivation for unexplored policies with the potential to be both efficiency enhancing and pro-poor, these heterogeneity findings also supply insights into other *pooled* elasticity estimates that do not consider heterogeneity.

While the natural-gas demand elasticity literature is sparse relative to that of the electricity literature [Rehdanz 2007], several previous papers offer estimates for the price elasticity of demand for residential natural gas. Table 1 lists the past studies that we found, the type of data used, and the resulting estimates of the own-price elasticity of demand. As Table 1 shows, past papers either estimate the elasticity of demand for residential natural gas using aggregated data (e.g., Hausman and Kellogg; Davis and Muehlegger) or using micro data with average prices (e.g., Alberini, Gans, and Velez-Lopez 2011; Meier and Rehdanz 2010). The exception is Rehdanz, who uses a two-period sample from West Germany, where it appears average price equalled marginal price. The majority of these papers do not attempt to deal with bias resulting from multiple sources of simultaneity, which we discuss below.

Research on the price elasticity of demand for natural gas faces two major challenges: insufficient data and multiple potential sources of endogeneity. Many of the available datasets aggregate consumption across both space and time. This aggregation—coupled with utilities' multi-tiered volumetric pricing regimes, income-based discounts, and fixed charges—makes it impossible for researchers to match consumers to the actual prices they face. Aggregation across customers and seasons also inhibits research into heterogeneity across consumers. Perhaps most importantly, research on the elasticity of demand for natural gas must also consider multiple potential sources of endogeneity. The first source of endogeneity is the classic simultaneity that stems from the fact that quantity and price result from the equilibrium in a system of equations. This is due to the fact that unlike the electricity sector, rates customers pay change on a monthly basis as a function of gas wholesale prices. The second

source of endogeneity results from the fact that price is mechanically a function of quantity in a block-rate price regime. As a household’s consumption increases, its marginal price also increases (consequently, its average price also increases).

2 Institutional setting

In this section we describe the basic organization of the natural gas industry, breaking the market into four segments: (1) production and processing, (2) transportation, (3) storage, and (4) local distribution companies (LDCs). Figure 1 illustrates the basic institutional organization of the natural gas industry.³ The four segments we discuss below roughly follow Figure 1 except that they exclude end users—those users who simply consume natural gas—and the liquid natural gas import/export-based segments of the market. While this paper focuses on the behavior of residential natural gas consumers, part of our identification strategy relies upon a basic understanding of the greater industry—specifically in understanding which instruments may shift supply without affecting demand. After discussing these four segments, we then describe the multi-tier pricing structure employed by the two Californian natural gas utilities in this paper.

2.1 Market segments

Production and processing Natural gas enters the market at the wellhead where it is produced and first sold [Brown and Yücel 1993]. Some wells produce only natural gas, while other wells produce natural gas in addition to crude oil [Levine, Carpenter, and Thapa 2014]. The raw product then moves from wellheads to processors. Processors remove impurities and separate the raw product into multiple commodities (separating “natural gas” from “natural gas liquids”) [Levine, Carpenter, and Thapa 2014].

Transportation High-pressure pipelines transport processed natural gas from production and processing areas to both intermediate users (storage facilities, processors, LDCs) and

³We include liquid natural gas (LNG) in the figure for completeness, but liquid natural gas does not play a large role in the natural gas market in the United States—LNG imports currently account for less than one percent of natural gas imports and accounted for three-percent of imports at their peak in 2007 [Levine, Carpenter, and Thapa 2014]. For this reason, we omit LNG for the rest of this paper.

final users (electricity generators, industrial users, commercial users, and residential users). Figure 2 maps this pipeline network for the continental United States. Private companies own and operate segments of the pipeline; these pipeline companies' rates are regulated at the state level and the national level [Levine, Carpenter, and Thapa 2014]. Extensive spot markets and futures markets exist for natural gas. Louisiana's Henry Hub connects to 13 intrastate and interstate pipelines. The Henry Hub is the designated delivery point for the New York Mercantile Exchange's natural gas futures contracts and the Henry Hub price is generally regarded as a nationally relevant price [Levine, Carpenter, and Thapa 2014]. Figure 3 depicts the Henry Hub spot price from 1997 through 2016. Transportation costs represent a substantial percentage of natural gas prices: according to Levine, Carpenter, and Thapa, in 2011–2012, 72 percent of consumers' average heating costs originated in "transmission and distribution charges".⁴ Thus, the natural gas transportation network creates a nationally integrated market and simultaneously contributes to a sizable portion of natural gas end-users' prices.

Storage Storage plays a major role in several parts the natural gas market, but all parties store mainly for the same reason: volatility within the market. Due to its major roles in heating and electricity production, natural gas demand is strongly driven by weather and can be unpredictable in the short run. To combat price volatility and to be able to meet peak demand, both local distribution companies and large natural gas consumers store gas underground [Levine, Carpenter, and Thapa 2014]. Producers utilize storage to smooth production.

Local distribution companies Local distribution companies' primary function is distributing natural gas to their contracted end users—industrial, residential, and commercial consumers of natural gas. To accomplish this task, LDCs purchase natural gas through both spot markets and longer-term contracts. In addition, LDCs own and operate their own pipeline and storage networks. To cover the fixed costs involved in their pipelines, storage, and ad-

⁴Levine, Carpenter, and Thapa also note that in 2007–2008 "transmission and distribution charges" accounted for 41 percent of consumers' average heating costs. It is worth keeping in mind that consumers' average heating costs fell approximately 20 percent in this period.

ministration, LDCs often utilize a combination of two-part tariffs and multi-tiered pricing regimes—though some utilities fold all of their costs into their volumetric pricing. State utility commissions (e.g., the California Public Utilities Commission) regulate LDCs' price regimes, allowing the LDCs to earn a regulated rate of return [Brown and Yücel 1993; Davis and Muehlegger 2010; Levine, Carpenter, and Thapa 2014].

2.2 Natural gas pricing in California

The California Public Utilities Commission (CPUC) regulates both utilities from whom we draw data in this paper—Pacific Gas and Electric Company (PG&E) and Southern California Gas Company (SoCalGas). Because this paper analyzes residential natural gas consumers' responses to natural-gas retail prices, the most relevant regulations facing PG&E and SoCalGas are CPUC's price and quantity regulations. In addition, the California Energy Commission (CEC) defines geographic climate zones (see Figure A1), which, in part, determine households' price schedules [California Energy Commission 2015, 2017].

For PG&E's and SoCalGas's residential consumers, a household's bill depends upon five elements:⁵

1. The two-tiered price schedule set by the utility
2. The total volume of natural gas consumed during the billing period
3. The season (*summer* or *winter*) in which the bill occurs
4. The climate zone into which the household's physical location falls
5. The household's CARE (California Alternate Rates for Energy) status

Figure A2 provides an example of a typical residential natural gas bill from PG&E.

Both PG&E and SoCalGas utilize two-tiered pricing regimes. The California Energy Commission (CEC) divides California into 16 climate zones in which households' needs for heating should be relatively homogeneous [California Energy Commission 2015, 2017; Pacific Gas

⁵Consumers' billing periods do not perfectly align with calendar months. However, PG&E's and SoCalGas's price changes do align with calendar months (during the years that our data cover). The two utilities deal with this misalignment of billing periods and price regimes slightly differently. PG&E calculates individual bills for each calendar month under the assumption that consumption is constant throughout the billing period. SoCalGas calculates a single bill using time-weighted average prices (averaging across the different price regimes). These methods are equivalent under a single linear price but differ under the actual multi-tiered price regimes. Please see the *Calculating bills* section in the appendix for more detail.

and Electric Company 2016]. The utilities also divide the year into heating (winter) and non-heating (summer) seasons. Based upon a household's climate zone (determined by the household's location) and the season, the CPUC determines a volume of natural gas that should be adequate for heating during the course of one day. This volume of natural gas is called the household's *daily allowance*. Multiplying the household's *daily allowance* by the number of days in the billing period gives the household's *total allowance* for the bill. For each unit (therm) of natural gas up to the bill's *total allowance*, the household pays the first tier's per-unit price (*baseline price*). For each unit of gas above the household's *total allowance*, the household pays the second tier's per-unit price (*excess price*).⁶ Figure 4a illustrates an example of the two-tier block-pricing regime used by PG&E and SoCalGas. Figure 4b depicts how residential consumers' (daily) tier-one allowances vary through time within a given climate zone (PG&E's climate zone R and SoCalGas's climate zone 1). Figure A1 depicts California's 16 California Energy Commission (CEC) defined climate zones.

Each month, the utilities update their price schedules. The absolute difference between the first-tier price and the second-tier price also varies but tends to remain constant for several months.⁷ These monthly price changes allow the utilities to charge customers at rates that reflect the prevailing price of natural gas—the utilities' costs—and are in fact tied to the utilities' costs, which directly correspond to spot market prices.⁸ If the utilities wish to change the way in which their prices are tied to market prices and other costs, they must receive authorization following a review process with CPUC. Figure 6a illustrates these monthly price-regime changes and the fairly fixed step between the two tiers. Figure 6b depicts the correlation between the utilities' baseline (first-tier) prices and the spot market price of natural gas at the Henry Hub.

A household's CARE (California Alternate Rates for Energy) status also affects the prices that the household faces. Households qualify for CARE by either meeting low-income qualifications or by receiving benefits from one of several state or federal assistance programs (e.g., Medi-Cal or the National School Lunch Program) [Southern California Gas Company

⁶The utilities in this paper work in units of volume called *therms*. One therm is equal to 100,000 Btu [U.S. Energy Information Administration 2016d].

⁷The utilities differ in their frequencies which they change this absolute difference: PG&E adjusts the distance between the two tiers' price much more frequently than SoCalGas.

⁸The utilities report their *weighted average costs of gas* to the CPUC.

[2016](#)]. CARE prices are 80 percent of standard prices at both tiers. In addition to giving us the household's correct pricing regime, we use CARE status to identify low-income households.

3 Data

3.1 Natural gas billing data

The billing data in this paper come from two major utilities in California: Pacific Gas and Electric Company (PG&E) and Southern California Gas Company (SoCalGas). The PG&E data cover residential natural gas bills in PG&E's territory from January 2003 through December 2014. The SoCalGas data cover residential natural gas bills from May 2010 through September 2015. Thus, the two utilities' data overlap from May 2010 through December 2014. After excluding zip codes with fewer than 50 households, PG&E's service area covers 597 5-digit zip codes (680,846 9-digit zip codes) with a total of 5,888,276 households and 180,663,705 bills. After excluding zip codes with fewer than 50 households, SoCalGas's service area covers 611 5-digit zip codes (610,207 9-digit zip codes) with a total of 2,526,503 households and 95,335,393 bills. The left side of Figure 5 depicts PG&E's and SoCalGas's service areas at the 5-digit zip code level. Table 2 provides a brief summary of the billing data with regards to the numbers of bills, households, zip codes, and monetary values of the bills. Tables 2 and 3 summarize prices, quantities, and other variables of interest—pooling across all observations and also splitting the data by season or CARE status. Both tables summarize the full dataset—all zip codes across both utilities—and a subset of the data based upon all 5-digit zip codes served by both utilities. We describe this subset in detail below in the *Empirical strategy* section.

The utilities' billing data are at the household-bill level: a single row of the dataset represents a single billing period for a given household. Table A16 describes the variables (columns) in this dataset. We follow the natural gas utilities' convention in defining a household (or customer) as the interaction between a unique utility account and a unique physical location identifier.

We also utilize historical data on pricing from the two utilities. As described above, these pricing data include (1) each utility's monthly two-tier pricing regime and (2) the daily

allowance for each climate zone during each season. After joining these pricing data to the households' billing data, we are able to determine both the marginal price and average price (and average marginal price) for each bill received by each household. Because the billing data include most households' 9-digit zip codes, we are able to match households to other datasets—for instance, Census block-group data.

3.2 Weather data

Data on daily weather observations originate from the PRISM project at Oregon State University [PRISM Climate Group 2004]. We match this local, daily weather data to the household consumption data at the day by zip-code level. This dataset contains daily gridded maximum and minimum temperature for the continental United States at a grid cell resolution of roughly 2.5 miles. We observe these daily data for California from 1980–2015. In order to match the weather grids to zip codes, we obtained a GIS layer of zip codes from ESRI, which is based on the US Postal Service delivery routes for 2013. For small zip codes not identified by the shape file we have purchased the location of these zip codes from a private vendor⁹. We matched the PRISM grids to the zip code shapes and averaged the daily temperature data across the multiple grids within each zip code for each day. For zip codes identified as a point, we simply use the daily weather observation in the grid at that point. This results in a complete daily record of minimum and maximum temperature—as well as precipitation—at the zip-code level from 1980–2015.

4 Empirical strategy

In this section we describe the empirical strategy we use to identify the price elasticity of demand for residential natural gas consumers. First, we present the basic estimating equation that drives the paper's results. Next, we discuss the inherent challenges to identification in this setting. We then discuss potential solutions to these challenges and detail which of these solutions are feasible in this paper's specific setting. Finally, before moving to the results, we provide evidence for the validity of the instruments.

⁹zip-codes.com

4.1 Estimating equation

The relationship at the heart of this paper's elasticity estimates is

$$\log(q_{i,t}) = \eta \log(p_{i,t}) + \lambda_{i,t} + \varepsilon_{i,t} \quad (1)$$

where i and t index household and time; q denotes quantity demanded; and p denotes price.¹⁰ The term $\lambda_{i,t}$ represent household fixed effects, time-based fixed effects, and/or household-by-time fixed effects—depending on the specification. Our main specification in this paper uses household by month-of-year fixed effects (e.g., household #1 in January) and zip-code by month-of-sample fixed effects (e.g., Fresno in January 2009; also called zip-code by year by month). A causally identified estimated of η yields the own-price elasticity of demand.

4.2 Challenges

Two main sources of endogeneity threaten identification in equation 1.

The first challenge in identifying this price elasticity of demand is the potential endogeneity that results from price and quantity being simultaneously determined by the equilibrium of supply and demand—*simultaneity* (e.g., Woolridge 2009). Standard ordinary least squares (OLS) fails to properly treat the endogeneity inherent in (1). As discussed above, many papers in the natural gas literature ignore this potential source of bias while estimating the price elasticity of demand—relying upon fixed effects, uncorrelated demand and supply shocks, and/or assumptions of exogenous prices. If simultaneity is indeed present in this setting, then the estimates in these papers will recover biased estimates for the elasticity of demand for residential natural gas.

A second challenge to identification in this paper results from our paper's specific context: the two-tiered pricing within the Californian natural gas market. Put simply, in tiered pricing regimes, the marginal price is a (weakly increasing, monotone) function of quantity. For the same reason, average price is also a function of quantity. Thus, when a household consumes more, its marginal and average prices mechanically increase. In terms of identifying the

¹⁰We first consider the price that classical economic theory deems relevant: the current period's marginal price. We also provide results using other measures of price, e.g., average price, average marginal price, and baseline (first-tier) price.

price elasticity of demand, this is *bad variation*: the marginal price that a household faces is endogenous because the marginal price is correlated with unobserved demand shocks [Ito 2014]. This bias is another form of simultaneity often called *reverse causality*.

In practice, one generally cannot sign the bias resulting from the classical simultaneity of price and quantity without making further assumptions regarding the correlation of supply and demand shocks. On the other hand, the bias resulting from marginal and average prices being a function of quantity results in upwardly biased estimates of demand *elasticities*. In extreme cases, this latter case of bias can yield estimates that suggest upward-sloping demand curves.

Table 4 demonstrates the consequences of failing to address these challenges to identification by estimating the price elasticity of demand— η in equation 1 via ordinary least squares (OLS) using marginal price (columns 1–3) and baseline (first-tier) price (columns 4–6). We also vary the set of controls for each price. For a given price, the leftmost columns utilize the simplest set of controls. The “identification strategy” present in Table 4 makes no attempt to correct for the aforementioned potential biases outside of a fairly rich set of fixed effects—a household by month-of-year fixed effect and a city by month-of-sample fixed effect. We also vary whether we control for heating degree days (HDDs) during the billing period. The rightmost columns for each price control for city by month-of-sample fixed effects.

The six regressions in Table 4 utilize two different measures of price: (1) the household’s marginal price during a given bill, and (2) the household’s baseline (first-tier) price during the given bill. These two—rather related—measures of price yield considerably different results—differing both quantitatively and qualitatively. The baseline price suggests an elasticity of approximately -0.10 , while the marginal price indicates a *positive* demand elasticity of approximately 1.05 . The substantial differences across estimates in Table 4 suggests at least one of the aforementioned biases are present. Specifically, the fact that the marginal-price based elasticity estimates are positive (implying upward-sloping demand curves), while the baseline-price based estimates are negative suggests the *price-is-a-function-of-quantity* flavor of simultaneity is a first-order problem in this context. This interpretation follows from the results due to the fact that baseline prices are not a function of quantity, while marginal prices are a function of quantity.

While the baseline price elasticity might appear to be reasonable in terms of magnitude, it is still not identified, as it still may suffer from simultaneity bias. Simply adding more observations in the flavor of the big data movement does not address this potential endogeneity.¹¹ The OLS results in Table 4 do not provide any of the obvious signs of the existence of the second type of simultaneity bias discussed above—classical simultaneity from price and quantity’s simultaneous determination in a system’s equilibrium. In the presence of such endogeneity, one might expect the inclusion of heating degree days or city by month-of-sample to alter the coefficients more than we observe in Table 4. However, this lack of considerable change in coefficients is not evidence of a lack of simultaneity. It is fundamentally a statistically untestable issue which stems from the theoretical setup of how market prices originate.

Finally, it is worth noting that the baseline-price based elasticity estimates are well within the range of estimates from the existing literature, as shown in Table 1. This outcome warrants some concern, as it suggests that some of these estimates may suffer from endogeneity.

4.3 Solutions

Having shown that OLS with fixed effects does not cleanly identify the own-price elasticity of demand in this setting, we now discuss several potential routes for identifying the causal effect of price on quantity in our setting. In the end, we opt for an identification strategy that combines a spatial discontinuity with an instrumental variables approach.

4.3.1 Discontinuities

A common route toward identification involves finding relatively small geographic units that receive different prices within the same time period. Arbitrary administrative boundaries that determine policies’ catchment areas provide a popular tool in this context, *e.g.*, Dell; Chen et al.; Ito. In our context of natural gas in California, the boundary between PG&E and SoCalGas offers potentially arbitrary within-city (and within-zip code) variation in prices during a month. Specifically, the boundary PG&E and SoCalGas bisects eleven cities—three

¹¹When we use this specification (with baseline price) in the full sample, we obtain estimates for the price elasticity of -0.33, -0.32, and -0.78 (specification/order corresponding to the three baseline-price columns of Table 4).

clusters—in southern California. The left panel of Figure 5 displays the two utilities’ service areas (sufficiently covered in the datasets). The right panel of Figure 5 zooms in on the eleven cities (39 zip codes) that PG&E and SoCalGas both serve. Within these eleven cities, PG&E serves all 39 zip codes, while SoCalGas serves 18 of the zip codes.

This identification strategy rests upon the assumption that households on one side of the utilities’ border provide a valid control group for the households on the other side of the border. Because the boundary mainly represents the extent of each utilities’ underground distribution network and is unlikely to enter into households’ preferences, the exogeneity of the boundary to household characteristics should be valid Ito. The main threat to this identification strategy is that utilities’ networks correlate with geographic or neighborhood characteristics over which individuals have preferences. However, we use household-month fixed effects, which absorb mean differences across households in a given calendar month. Thus, for the border discontinuity to be invalid, households must sort in a way consistent with their elasticities, and the utilities’ price series must differ significantly in their variances. Because the data contain considerable variation in prices for both utilities and the panel contains approximately six years of monthly bills, this sort of sorting bias seems unlikely. Figure 6b suggests the generating distributions for the utilities’ prices are quite similar (the standard deviation of the price series are 0.0940 and 0.1053 for PG&E and SoCalGas, respectively).

Ito employs a similar strategy within the context of electricity consumption. However, there is at least one significant difference between the electricity and natural gas contexts which prevent us from completely adopting Ito’s identification strategy: discontinuities within electricity utilities’ seven-tier pricing regime. By law, the electricity utilities in Ito’s study are not allowed to move the price of the first two tiers—they must recover changes in their costs by moving tiers three through seven. Thus, marginal prices in Ito’s setting move differently depending upon a household’s tier. Ito argues that the residual variation—combining the spatial discontinuity with this pricing discontinuity and geographic and temporal fixed effects—is plausibly exogenous from demand shocks. Because natural gas has only two tiers—and because the absolute difference between the two tiers has relatively low variation—we are unable to take advantage of price-tier based discontinuities. Therefore, in addition to this utility-border-based discontinuity, we adopt an additional strategy to overcome endogeneity.

4.3.2 Instrumental variables

The second element of our estimation strategy for identifying the price elasticity of demand for natural gas involves a traditional solution to simultaneity: supply-shifting instruments. In this context, the ideal supply-shifting instrument is (1) strongly correlated with the prices that the natural gas utilities charge their customers (the *first stage*), and (2) uncorrelated with residual shocks affecting consumers' demand [Angrist and Pischke 2009]. In this paper, our instrument is the Henry Hub spot price for natural gas.

Henry Hub spot price Specifically, we instrument the prices that consumers face (e.g., marginal price, average price, baseline price) with the average spot price at Louisiana's Henry Hub in the weeks preceding the change in prices. We also interact the spot price with *utility* to allow the utilities to differentially incorporate price changes. The Henry Hub spot price represents the nationally prevailing price for short-term natural gas contracts (the hub sits at the intersection of 13 intrastate and interstate pipelines). This instrument satisfies the requirement of having a strong first stage, as the utilities base their prices, in part, on market prices of natural gas—the utilities buy natural gas on the spot market, and the California Public Utilities Commission regulates how the utilities fold their costs into the price regimes that customers face on a monthly basis.

The exclusion restriction for this spot-price based instrument is less obvious. However, several factors suggest that the exclusion restriction is plausibly valid. First, California's entire residential natural gas demand represents *at most* three-percent of national natural gas consumption—limiting the individual utilities' ability to set/influence spot prices and the Henry Hub. Second, we interact the spot price instrument with *utility*. This interaction, conditional on zip-code by month-of-sample fixed effects, implies that the identifying variation in our instruments comes from the difference across the ways the two utilities' incorporate monthly spot-price shocks into their pricing regimes. Third, because the utilities must obtain approval for price changes before the new price regime begins, the spot price is temporally disconnected from the billing period. In other word, the utilities' costs (and approved prices) are based upon spot prices that precede the billing period by several weeks. Thus, shocks

that affect the Henry Hub spot price are distinct in time from shocks that affect natural gas demand—our fixed effects will absorb any of these shocks, so long as they do not differ across the utilities’ border within a month. Finally, we control for the number of heating degree days (HDDs) in the household’s zip code, during the households’ billing period. Because residential consumers primarily use natural gas in heating applications, controlling for HDDs further reduces the opportunity for local demand shocks affecting national. Thus, we argue that the exclusion restriction is plausibly valid for our spot-price instrument.

Spot price first stage Panel (A) of Table 5 provides the first-stage estimates for the two-stage least squares equations

$$\log(p_{i,t}) = \pi_{1a} p_{i,t}^{\text{spot}} + \pi_{1b} p_{i,t}^{\text{spot}} \times \text{SCG}_i + \pi_2 \text{HDD}_{i,t}^{\text{bill}} + \text{HH}_{i,t} + \text{Zip}_{i,t} + u_{i,t} \quad (2)$$

$$\log(q_{i,t}) = \eta_1 \widehat{\log(p_{i,t})} + \eta_2 \text{HDD}_{i,t}^{\text{bill}} + \text{HH}_{i,t} + \text{Zip}_{i,t} + v_{i,t} \quad (3)$$

where $\text{HH}_{i,t}$ is a household by month-of-year fixed effect, $\text{Zip}_{i,t}$ is a zip-code by month-of-sample fixed effect, and SCG_i is an indicator for whether the household’s retail utility is SoCalGas. Figure 6b gives visual first-stage evidence—illustrating the link between the two utilities’ prices and the Henry Hub spot price and also demonstrating how the utilities differ in their response to the spot price.

Panel (A) of Table 5 gives the first-stage results corresponding to equation 2. Specifically, the Henry Hub spot price shown in Table 5 is the average natural gas spot price at Henry Hub during the 7 days preceding the change in pricing. Table 5, Panel A displays the first-stage results for five different prices that may be relevant to households: marginal price, average price, average marginal price, baseline price, and simulated marginal price¹² (using the log of each variable). Each price is the second lag of the contemporaneous price—the current bill is lag zero, the prior bill contains the first lags of price, and the bill preceding the prior bill contains the second lags of price. Tables A3–A7 illustrate consumers’ varying responses to different lags of price.

Both Figure 6b and Panel A of Table 5 demonstrate that the spot-price based instruments

¹²Simulated marginal price refers to a simulated instrument for marginal price. We discuss this measure of price in the next section.

are quite strong (the F statistics testing the joint significance of the instruments exceed 350 in all cases). This significance is unsurprising, as the utilities purchase gas on the spot market and incorporate these costs directly into their price regimes. The significance of the interaction between spot price and utility (SoCalGas) in Table 5 implies the utilities differ appreciably in the ways that they incorporate spot-market costs into their pricing regimes—PG&E’s pricing regime appears to be much less responsive to the spot price than SoCalGas.¹³ Though the city by month-of-sample fixed effect should control for most local demand shocks, bills do not perfectly match months. The within-bill HDDs variable controls for any remaining weather-based demand shocks. The results in Table A10 demonstrate robustness to including within-bill heating degree days (the even-numbered columns), which suggests that the instrument is exogenous to local-California weather shocks, one of the key local-demand drivers in natural gas [Davis and Muehlegger 2010; Levine, Carpenter, and Thapa 2014; Hausman and Kellogg 2015].

While the first stage is quite strong for all specifications, the results in Panel A of Table 5 suggest the instrument is strongest—in terms of first-stage significance—for baseline price, followed by average marginal price, average price, marginal price, and finally simulated marginal price. A likely reason for this outcome is that baseline price is the least noisy price: it is the only price that is not a function of the consumer’s quantity, and it does not include variation from changes in the size of the step between the two tiers’ prices. By these terms, (simulated) marginal price, is the noisiest, which is consistent with marginal price having the small first-stage F statistic of three prices.

4.3.3 Instrumented prices and simulated instruments

In the preceding sections, we discussed how our spot-price based instrument—used in conjunction with the spatial discontinuity in utilities’ service areas—may overcome the bias resulting from the fact that quantity and price (our dependent and independent variables) result from a simultaneously determined equilibrium. We now discuss the aspect our identification strategy that deals with the *price-is-a-function-of-quantity* endogeneity present in multi-tiered pricing

¹³One difference between the utilities’ pricing regimes is that PG&E does not have a fixed charge, while SoCalGas does. Thus, PG&E recovers both fixed and volumetric costs through volumetric charges to its customers.

contexts.¹⁴ We present three separate options for breaking this endogenous link between price and quantity, but in the end, the options yield very similar results.

Option 1: Instrumented prices One method for breaking the endogenous link between a household's price and its quantity is simply to instrument the household's price with a variable that is aggregated at a unit above household. Consider the IV strategy discussed above: instrumenting a household's price with the our Henry Hub spot price interacted with utility. Because this instrument only varies at the billing-period by utility level, when we regress a household's endogenous price on this instrument (and our set of fixed effects) in the first stage, the variation captured by the predicted prices is only the variation that correlates with the spot price, which is determined weeks before the household's consumption decision. Thus, if the spot price provides a valid instrument for the classical simultaneity context, it also provides a valid instrument for the second *price-is-a-function-of-quantity* endogeneity.

Option 2: Baseline price In a similar manner, the baseline prices provides a valid instrument that breaks the *price-is-a-function-of-quantity* endogeneity. Because a household's baseline price is not a function of its quantity consumed, baseline price does not suffer from the same endogeneity. Baseline price is also strongly predictive of marginal (or average) price. Thus, in application, one could either replace marginal (or average) price with baseline price or instrument one of the endogenous prices with baseline price. However, baseline price fails to capture the higher price that a household faces once it exceeds its total monthly allowance.

Option 3: Simulated instrument Simulated instruments provide a third option for breaking the *price-is-a-function-of-quantity* flavor of endogeneity. The simulated-instrument approach follows a methodology suggested by Ito. Specifically, this approach creates an instrument (or proxy) for marginal (or average) price by plugging a lagged level of consumption into the current price regime, *i.e.*,

$$z_{i,t} = p_{i,t}(q_{t-k}) \quad (4)$$

¹⁴This endogeneity is present both in marginal price and in average price.

The main idea for this instrument is using a household's consumption history to predict whether a household will face the baseline or excess price in the current period. As with any instrument, we want to accomplish this prediction in a way that is (1) strongly predictive of the true outcome (the first stage) and that is uncorrelated with any recent shocks to the household (the exclusion restriction) [Angrist and Pischke 2009]. For these reasons, we modify the equation 4 slightly. First, we use the households' lagged consumption levels (from lagged bills 10 through 14 months prior) to determine in how many of the lagged periods would exceed this billing period's baseline allowance, *i.e.*,

$$v_{i,t} = \frac{1}{5} \sum_{k=10}^{14} \mathbb{1}\{q_{i,t-k} > \bar{A}_{i,t}\} \quad (5)$$

where \bar{A}_{it} is household i 's baseline allowance in time t . We then calculate the *simulated instrument* for marginal price, $z_{i,t}$ via

$$z_{i,t} = \mathbb{1}\{v_{i,t} \leq 0.5\} \times p_{i,t}^{\text{base}} + \mathbb{1}\{v_{i,t} > 0.5\} \times p_{i,t}^{\text{excess}} \quad (6)$$

Summarizing equations 5 and 6: this simulated instrument predicts that a household will exceed its allowance when the household exceeded its allowance in the majority of past bills (based upon lagged months 10 through 14).¹⁵

Column (5) of Table 5 (Panel A) provides the first-stage results consistent with equation 2 but with the simulated instrument of marginal price substituted for actual marginal price (we still instrument with the Henry Hub spot price).¹⁶ The first-stage is again quite strong in this specification, and the results are qualitatively similar to the results in columns (1)–(4) of Table 5, Panel A. Henceforth we will refer to the simulated instrument for marginal price as *simulated marginal price*.

For the sake of simplicity, all subsequent results jointly employ both the spatial discontinuity and the spot-price based instrument. To incorporate the competing strategies discussed

¹⁵This simulated instrument is robust to the choice of months 10 through 14. The goal is to keep the instrument in the same season as the current bill (keeping the first stage strong), while allowing some distance from the current period (the exclusion restriction: preventing medium-run shocks from affecting both periods).

¹⁶It is worth noting that in this paper, any result utilizing the simulated instrument will have fewer observations than other results, as the simulated instrument is greedier for data—for an observation to remain in the dataset, its 14th lag must also be in the dataset.

immediately above, we provide results consistent with the strategies: instrumenting with spot price, utilizing baseline price, and utilizing *simulated* marginal price (the simulated instrument for marginal price). We now turn to our main results.

5 Results

In this section, we discuss the estimated price elasticities, using the empirical strategies extensively discussed above. After presenting the main results for the *pooled* elasticity (no heterogeneity), we further examine whether households' price responses (*i.e.*, elasticities) vary by season and/or by income.

5.1 Pooled price elasticity of demand for natural gas

Panel B of Table 5 displays the elasticity results from the second-stage regression in equation 3. These results instrument log price with the Henry Hub spot price (interacted with utility), exploit the spatial discontinuity at the zip-code level, and use the log of daily average consumption (in therms) as the outcome. The five columns each estimate the elasticity using the log of different type of price: marginal price, average price, average marginal price,¹⁷ baseline price, and simulated marginal price. As discussed above, each price is the second lag of price, as opposed to the contemporaneous price. The estimated elasticity is robust to the type of price. Table A11 demonstrates that the estimated elasticity is also robust to the inclusion/exclusion of heating degree days or levels of fixed effects. This robustness to type of price also demonstrates robustness to how we control for the *price-is-a-function-of-quantity* endogeneity discussed above. The estimates for the price elasticity of demand range from -0.17 to -0.23.

Compared to their OLS-based counterparts in Table 4, the marginal-price based estimates for the elasticity of demand now have opposite—and theoretically correct—signs. The magnitudes of the estimates of the elasticity are theoretically reasonable and within the range of previous findings. Furthermore, these estimates, are plausibly identified and utilize

¹⁷We define *average marginal price* as the quantity-weighted marginal price paid by a customer during her billing period. *Average marginal price* does not include fixed charges, while *average price* does.

consumers' actual prices.

As discussed above, the results discussed up to this point—e.g., the results in Table 5—estimate the price elasticity of demand for residential natural gas using the second lag of the various prices. In order for a household to know the prices of its contemporaneous bill, the household would need to closely follow the approved advice letters published online by the utility or the California Public Utilities Commission. Because a household will not receive its previous bills for several days into its current billing period—and the household may not know view the previous bill until it pays the bill (or its credit card bill, if the household uses automatic bill payment) weeks later—the household may not know the prices from the previous bill until the current period is nearly over. For these reasons, it is plausible that the second lag of price is the most salient price to many households.

Table 6 replicates the second-stage results for marginal price and average price but varies the lag, ranging from contemporaneous price, to the first lag of price, to the second lag of price. Tables A3–A7 vary the lag of price, ranging from the household's contemporaneous price to the third lag of price. Across the five types of price, neither the contemporaneous price, nor the first lag of price, differ significantly from zero. For each type of price, both the second and third lags of price differ significantly from zero. For each price, the second-lag based elasticity slightly exceeds the third-lag based elasticity in magnitude, but the difference does not exceed the standard error. These results are consistent with households responding to two-to-three lags of price, as opposed to contemporaneous price.

Tables A3–A7 further examine the robustness of the elasticity estimates to the specification of spatiotemporal fixed effects and the inclusion/exclusion of within-bill heating degree days (HDDs). Each table examines the robustness of the elasticity estimates for a given type of price—marginal price (Table A3), simulated marginal price (Table A4), average marginal price (Table A5), average price (Table A6), and baseline (first-tier) price (Table A7). As shown in Panel (B) of Table 5, across five types of price, the point estimates for the price elasticity of demand for residential natural gas range from -0.17 to -0.23. Tables A3–A7 illustrate robustness of the estimated elasticities to specification.

5.2 Heterogeneity

We now examine the evidence that the price elasticity of demand for natural gas varies across income levels and/or seasons. If heterogeneity exists, then the regressions in the preceding section *pool* across the heterogeneous effects. This pooled parameter estimate may still be relevant for policy applications—particularly for policies that cannot differentiate between seasons or income groups. However, because OLS weights heterogeneous treatment effects by their share of the residual variation in the variable of interest—which is itself a function of (1) the numbers of observations in the heterogeneous groups and (2) the (residual) within-group variance in the variable of interest [Solon, Haider, and Wooldridge 2015]—one might wonder whether the pooled estimator always provides a policy relevant estimate. In addition, in the presence of heterogeneous elasticities, policymakers can increase efficiency by integrating these (known) heterogeneities [Ramsey 1927; Boiteux 1971; Davis and Muehlegger 2010].

For income-based heterogeneity, we use a household’s CARE status as a proxy for its income level.¹⁸ As discussed above, households qualify for CARE by either meeting low-income qualifications or by receiving benefits from one of several state or federal assistance programs (e.g., Medi-Cal or the National School Lunch Program) [Southern California Gas Company 2016]. For seasonal heterogeneity, we split the calendar into winter months (October through March) and summer months (April through September).¹⁹

5.2.1 Income heterogeneity

To examine income-based heterogeneity in the price elasticity of demand for natural gas, we estimate the two-stage least squares equations 2 and 3 separately for California Alternative Rates for Energy (CARE) households and non-CARE households. Columns (3) and (4) of Table 5 display the second-stage results from these regressions, providing estimates of the elasticity of demand by income level (CARE status).

¹⁸Because we do not have identifying variation in income level (or season), the heterogeneities that we estimate should be taken as descriptive statistics for the given group, rather than causal estimates. In other words, while we estimate heterogeneous elasticities with respect to income level, this heterogeneity may have nothing to do with income and could instead result from some other (omitted) variable that correlates with income/CARE status. However, identification of the sources of heterogeneity is not the goal of this paper: we aim to identify the elasticity of demand and demonstrate dimensions of heterogeneity. We leave it for future papers to identify the sources of these heterogeneities.

¹⁹This definition reflects southern California’s two seasons: warm and slightly less warm.

The results in columns (3) and (4) of Table 7 suggest that the elasticity results in the previous section may in fact pool across heterogenous elasticities: we estimate the price elasticity for CARE (lower-income) households is approximately twice that of non-CARE (higher-income) households. Specifically, we estimate an elasticity of approximately -0.24 (0.080) for CARE households and -0.14 (0.068) for non-CARE households. The “pooled” estimate corresponding to these results is -0.21 (0.071) (column (1) of Panel B in Table 5)—slightly higher than the midpoint between the CARE estimate than the non-CARE estimate.

5.2.2 Seasonal heterogeneity

To estimate seasonal heterogeneity in the price elasticity of demand for residential natural gas, we estimate the two-stage least squares equations 2 and 3 separately for winter months and for summer months. Columns (1) and (2) of Table 7 display the second-stage results from these regressions, providing estimates of the elasticity of demand by season. We again vary the inclusion of HDDs for robustness.

The results in Table 7 indicate a stark and significant difference between price elasticities in summer and winter months. The estimated price elasticity of demand for natural gas in summer months is approximately 0.052 (0.029) and marginally differs significantly from zero. The estimated elasticity for winter months is approximately -0.38 (0.14) and differs significantly from zero at the 1-percent level. The comparable “pooled” elasticity estimate corresponding to these results is approximately -0.21 (0.071) (column (1) of Panel B in Table 5). These results provide strong evidence that households’ consumptive and price-response behaviors vary considerably by season.²⁰

5.2.3 Income-by-season heterogeneity

Having shown potential heterogeneity across income groups (CARE status) and season, we now examine the evidence that income groups’ heterogeneity varies by season—interacting the heterogeneity dimensions discussion above (income and season).

To estimate seasonal heterogeneity in the price elasticity of demand for residential natural gas, we estimate the two-stage least squares equations 2 and 3 separately for winter months

²⁰Table A8 reproduces these heterogeneity results using average price, rather than marginal price.

and for summer months. Table 7 displays the second-stage results from these regressions, providing estimates of the elasticity of demand by season and CARE status. All results in Table 8 include HDDs.

The results in Table 8 are consistent with heterogeneous elasticities that depend upon the interaction between household income (CARE status) and season. In other words, the difference between a household’s winter and summer price elasticities varies by the household’s income level (CARE status). Specifically, the results in Table 8 indicate that both income groups are essentially inelastic to prices in summer months: we estimate a “summertime” price elasticity of 0.046 (0.035) for CARE households and 0.074 (0.032) for non-CARE households. Both elasticities are positive, but only one is significantly different from zero and small. In winter months, both sets of consumers are significantly more elastic and the elasticities are negative, but CARE households are especially more price elastic. We estimate the “wintertime” price elasticity of demand for natural gas is -0.523 (0.142) for CARE households and -0.317 (0.150) for non-CARE households. Again, the pooled elasticity corresponding to these results is approximately -0.19 (0.043) (column (2) in Table A3), which is a bit lower than the average of these four elasticities. The estimated effects of within-bill heating degree days in Table 8 also bear evidence of this two-way heterogeneity: the difference between responses to HDDs in summer and winter months appears to differ by income group. Overall, Table 8 demonstrates the potential for substantial and important heterogeneity underlying commonly estimated pooled elasticities.

6 Conclusion

This paper combines millions of household natural gas bills with manifold identification strategies to provide the first micro-data based causal estimates of the own-price elasticity of demand for residential natural gas. Utilizing cross-border price variation between California’s two largest natural gas utilities, in conjunction with a supply-shifting instrument that generates plausibly exogenous variation in price, the preferred specification results in an estimated elasticity of -0.21 $[-0.35, -0.07]$. This estimate is robust to specification choices which include within-bill weather, a number of different price instruments, and definition of

price. The point estimates for the own-price elasticity range from -0.23 to -0.17 for different measures of price. Given the robustness of these findings, this paper provides tight bounds on a policy-relevant parameter key to applications ranging from estimating the welfare benefits of fracking [Hausman and Kellogg 2015] to analyzing the regressivity of two-part tariffs [Borenstein and Davis 2012].

As a second important finding, we estimate that the own-price elasticity of demand varies significantly across seasons and customer types. We show that households on a popular low-income program, which subsidizes households' natural gas and electricity, appear to be twice as elastic in their response to price as households who are not part of the program. We also show that the price elasticity varies greatly across seasons. If we average across types of households, the summer price elasticity is close to, and only marginally different from, zero. The winter price elasticity is -0.38 . This heterogeneity suggests that households are much more price sensitive during their high-consumption months—the winter. These high-consumption winter months also correspond to the time of year in which consumers use natural gas in its most salient form: heating. When we break down the price elasticity across users and seasons, we show that subsidized consumers display the largest price sensitivity during the winter (-0.52). Neither type of customer displays a significant price response in the summer. These results suggest that if suppliers want to pass through costs to consumers, summertime is best—both for efficiency and for progressivity.

References

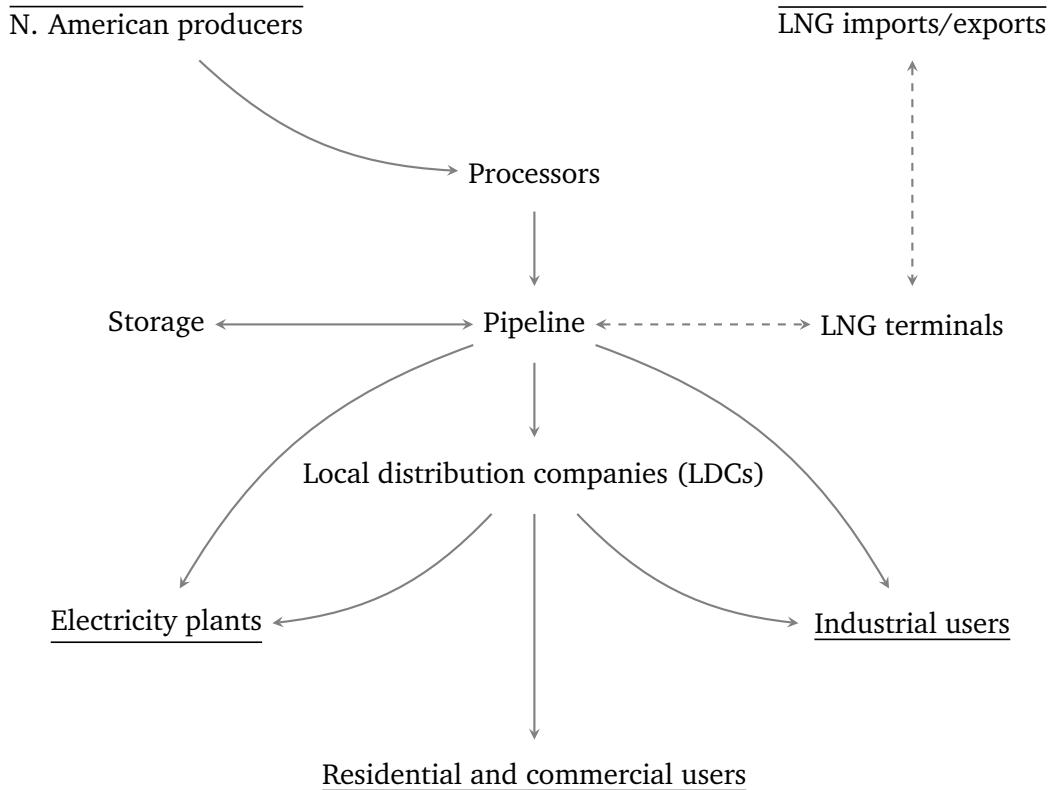
- Alberini, Anna, Will Gans, and Daniel Velez-Lopez. 2011. "Residential consumption of gas and electricity in the U.S.: The role of prices and income". *Energy Economics* 33 (5): 870–881. ISSN: 01409883. doi:[10.1016/j.eneco.2011.01.015](https://doi.org/10.1016/j.eneco.2011.01.015). <http://dx.doi.org/10.1016/j.eneco.2011.01.015>.
- Angrist, Joshua D., and Jörn-Steffen Pischke. 2009. *Mostly harmless econometrics: an empiricist's companion*. Princeton University Press.
- Boiteux, M. 1971. "On the management of public monopolies subject to budgetary constraints". *Journal of Economic Theory* 3 (3): 219–240. ISSN: 10957235. doi:[10.1016/0022-0531\(71\)90020-2](https://doi.org/10.1016/0022-0531(71)90020-2).
- Borenstein, Severin, and Lucas W. Davis. 2012. "The Equity and Efficiency of Two-Part Tariffs in U.S. Natural Gas Markets". *Journal of Law and Economics* 55.
- Brown, SPA, and MK Yücel. 1993. "The Pricing of Natural Gas in U.S. Markets". *Economic Review-Federal Reserve Bank of Dallas*, no. April: 41–52. <http://www.dallasfed.org/assets/documents/research/er/1993/er9302c.pdf>.
- California Energy Commission. 2015. "California Building Climate Zone Areas". Accessed March 01, 2017. http://www.energy.ca.gov/maps/renewable/building_climate_zones.html.
- . 2017. "Title 24: Building Energy Efficiency Program". Accessed March 01, 2017. <http://www.energy.ca.gov/title24/>.
- Chen, Yuyu, et al. 2013. "Evidence on the Impact of Sustained Exposure to Air Pollution on Life Expectancy from China 's Huai River Policy". *Proceedings of the National Academy of Sciences of the United States of America* 110 (32): 12936–12941. doi:[10.1073/pnas.1300018110/-/DCSupplemental.www.pnas.org/cgi/doi/10.1073/pnas.1300018110](https://doi.org/10.1073/pnas.1300018110/-/DCSupplemental.www.pnas.org/cgi/doi/10.1073/pnas.1300018110). <http://www.pnas.org/content/early/2013/07/03/1300018110.short>.
- Davis, Lucas W., and Erich Muehlegger. 2010. "Do Americans consume too little natural gas? An empirical test of marginal cost pricing". *RAND Journal of Economics*. ISSN: 07416261. doi:[10.1111/j.1756-2171.2010.00121.x](https://doi.org/10.1111/j.1756-2171.2010.00121.x).
- Dell, Melissa. 2010. "The Persistent Effects of Peru's Mining Mita". *Econometrica* 78 (6): 1863–1903. ISSN: 0012-9682. doi:[10.3982/ECTA8121](https://doi.org/10.3982/ECTA8121). <http://doi.wiley.com/10.3982/ECTA8121>.
- Environmental Protection Agency. 2013. "Natural Gas Extraction - Hydraulic Fracturing". Accessed August 05, 2016. <https://www.epa.gov/hydraulicfracturing>.
- Hausman, Catherine, and Ryan Kellogg. 2015. "Welfare and Distributional Implications of Shale Gas". *Brookings Papers on Economic Activity*, no. Spring: 71–125.

- Ito, Koichiro. 2014. "Do Consumers Respond to Marginal or Average Price? Evidence from Nonlinear Electricity Pricing". *American Economic Review* 104 (2): 537–563. ISSN: 0002-8282. doi:[10.1257/aer.104.2.537](https://doi.org/10.1257/aer.104.2.537). <http://pubs.aeaweb.org/doi/abs/10.1257/aer.104.2.537>.
- Levine, Steven, Paul Carpenter, and Anul Thapa. 2014. *Understanding natural gas markets*. Tech. rep. American Petroleum Institute. <http://www.api.org/oil-and-natural-gas/energy-primers/natural-gas-markets>.
- Meier, Helena, and Katrin Rehdanz. 2010. "Determinants of residential space heating expenditures in Great Britain". *Energy Economics* 32 (5): 949–959. ISSN: 01409883. doi:[10.1016/j.eneco.2009.11.008](https://doi.org/10.1016/j.eneco.2009.11.008). <http://dx.doi.org/10.1016/j.eneco.2009.11.008>.
- National Academy of Sciences. 2016. "What You Need to Know About Energy: Natural Gas". Accessed August 05, 2016. <http://needtoknow.nas.edu/energy/energy-sources/fossil-fuels/natural-gas/>.
- Pacific Gas and Electric Company. 2016. "Guide to California Climate Zones". Accessed August 13, 2016. <http://www.pge.com/myhome/edusafety/workshopstraining/pec/toolbox/arch/climate/index.shtml>.
- PRISM Climate Group. 2004. "PRISM Gridded Climate Data". Oregon State University. <http://prism.oregonstate.edu>.
- Ramsey, F. P. 1927. "A Contribution to the Theory of Taxation". *The Economic Journal* 37 (145): 47–61.
- Rehdanz, Katrin. 2007. "Determinants of residential space heating expenditures in Germany". *Energy Economics* 29 (2): 167–182. ISSN: 01409883. doi:[http://dx.doi.org/10.1016/j.eneco.2006.04.002](https://doi.org/10.1016/j.eneco.2006.04.002). <http://www.sciencedirect.com/science/article/pii/S0140988306000405>.
- Solon, Garry, Steven J. Haider, and Jeffrey Wooldridge. 2015. "What are we weighting for?" *Journal of Human Resources* 50 (2): 301–316. doi:[10.3386/jhr.50.2.301](https://doi.org/10.3386/jhr.50.2.301).
- Southern California Gas Company. 2016. "California Alternate Rates for Energy (CARE)". Accessed August 08, 2016. <https://www.socalgas.com/save-money-and-energy/assistance-programs/california-alternate-rates-for-energy>.
- U.S. Energy Information Administration. 2016a. "Average annual natural gas spot price in 2015 was at lowest level since 1999". Accessed November 04, 2017. <https://www.eia.gov/todayinenergy/detail.php?id=24412#>.
- . 2016b. "Electric Power Monthly - Net Generation by Energy Source: Total (All Sectors), 2006-May 2016". Accessed August 05, 2016. https://www.eia.gov/electricity/monthly/epm_table_grapher.cfm?t=epmt_1_1.

- . 2016c. “Energy-related CO₂ emissions from natural gas surpass coal as fuel use patterns change”. Accessed February 28, 2017. <https://www.eia.gov/todayinenergy/detail.php?id=27552>.
 - . 2016d. “Frequently Asked Questions”. Accessed March 01, 2017. <https://www.eia.gov/tools/faqs/faq.cfm?id=45&t=8>.
 - . 2016e. “Henry Hub Natural Gas Spot Price”. Accessed August 10, 2016. <https://www.eia.gov/dnav/ng/hist/rngwhhdD.htm>.
 - . 2017. “Maps: Layer Information for Interactive State Maps - Natural Gas Interstate and Intrastate Pipelines”. Accessed October 23, 2017. https://www.eia.gov/maps/layer_info.m.php.
- Woolridge, Jeffrey. 2009. *Introductory Econometrics - A Modern Approach*. 4e. Mason, OH: South-Western Cengage Learning. ISBN: 978-0-324-66054-8.

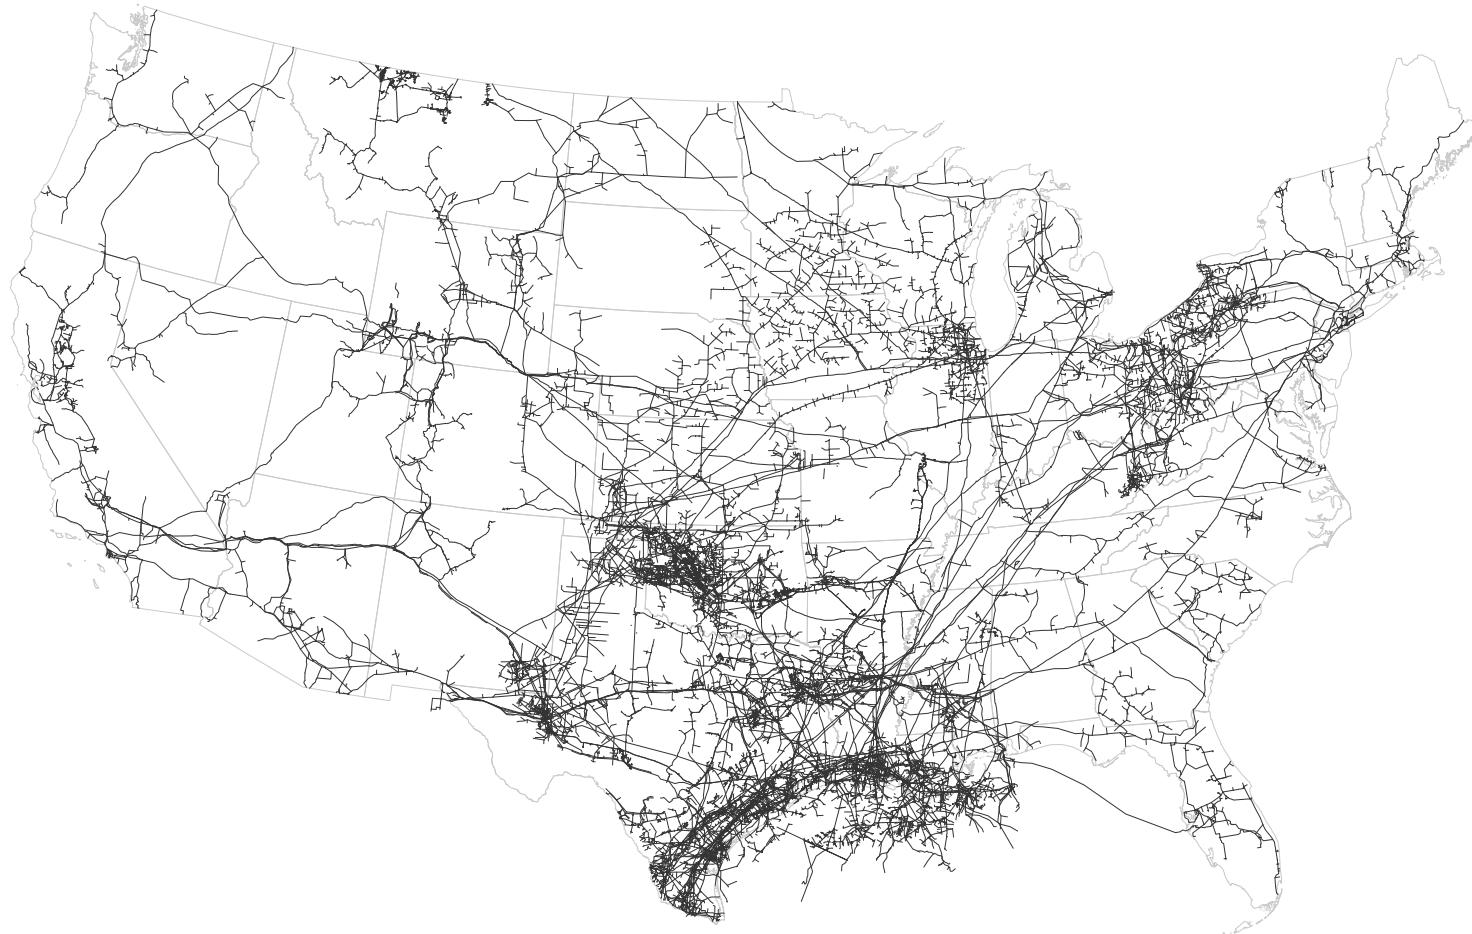
7 Figures

Figure 1: U.S. Natural gas institutional organization



Notes: Overbars represent points of entry into the U.S. natural gas market; underbars represent end points in the market; all other labels represent intermediaries. Arrow directions correspond to the direction of the flow of natural gas. The acronym LNG abbreviates *liquid natural gas*. This figure is based off of Levine, Carpenter, and Thapa with modification following Brown and Yücel.

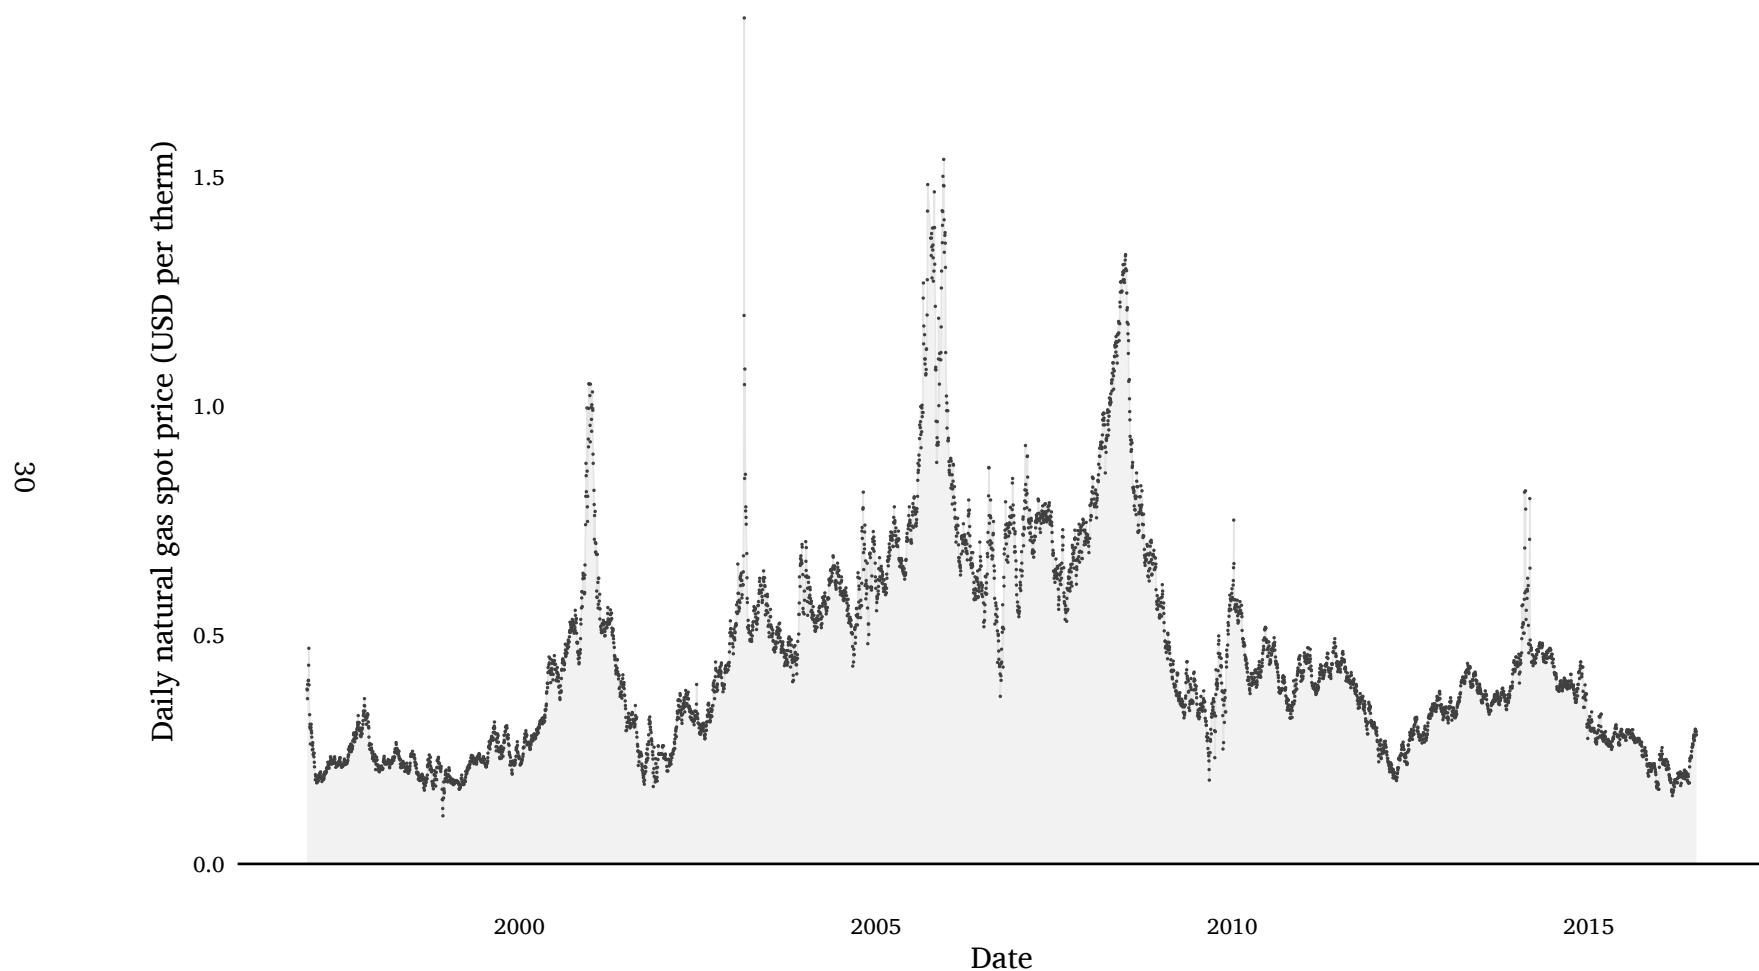
Figure 2: U.S. natural gas pipeline network



29

Notes: This map depicts the intra- and inter-state natural gas pipeline network for the (continental) United States (in black) overlayed on a map of the (continental) U.S. (light grey). *Source:* U.S. Energy Information Administration

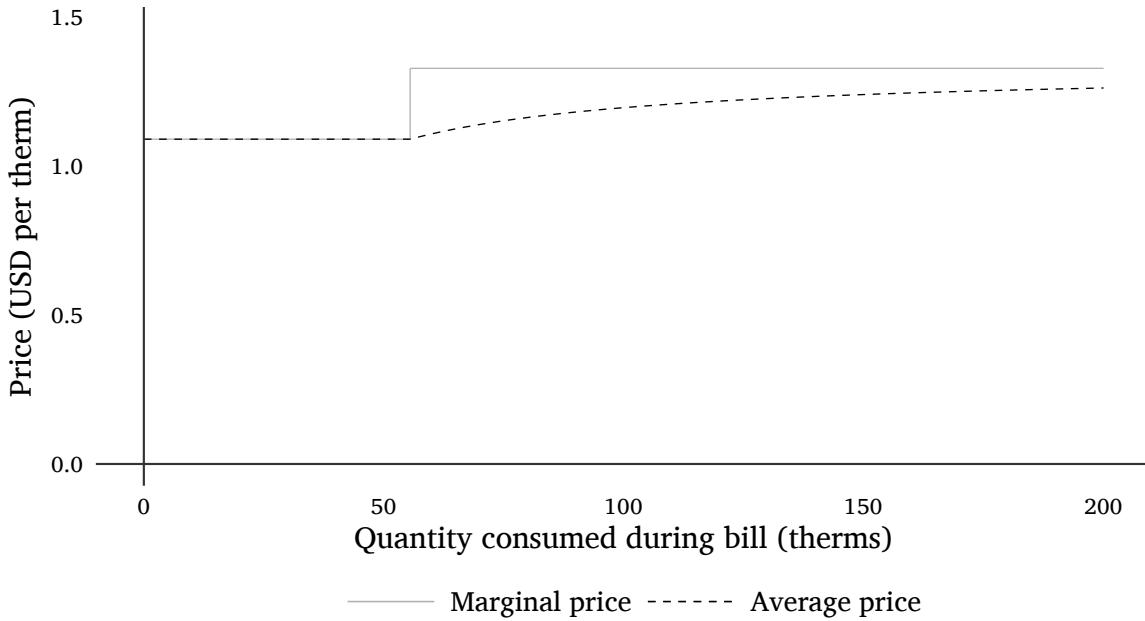
Figure 3: Henry Hub natural gas spot price: Daily, 1997–2016



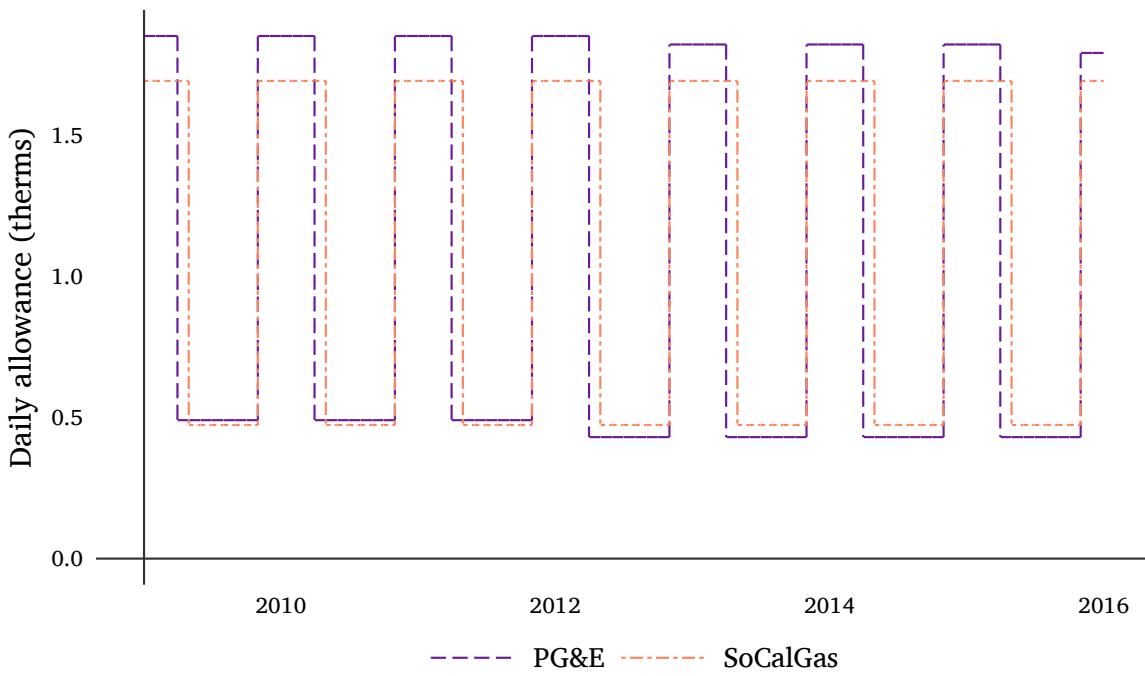
Source: U.S. Energy Information Administration

Figure 4: Households' allowances and prices

(a) Allowance and marginal vs. average price example: PG&E, January 2009, climate zone R



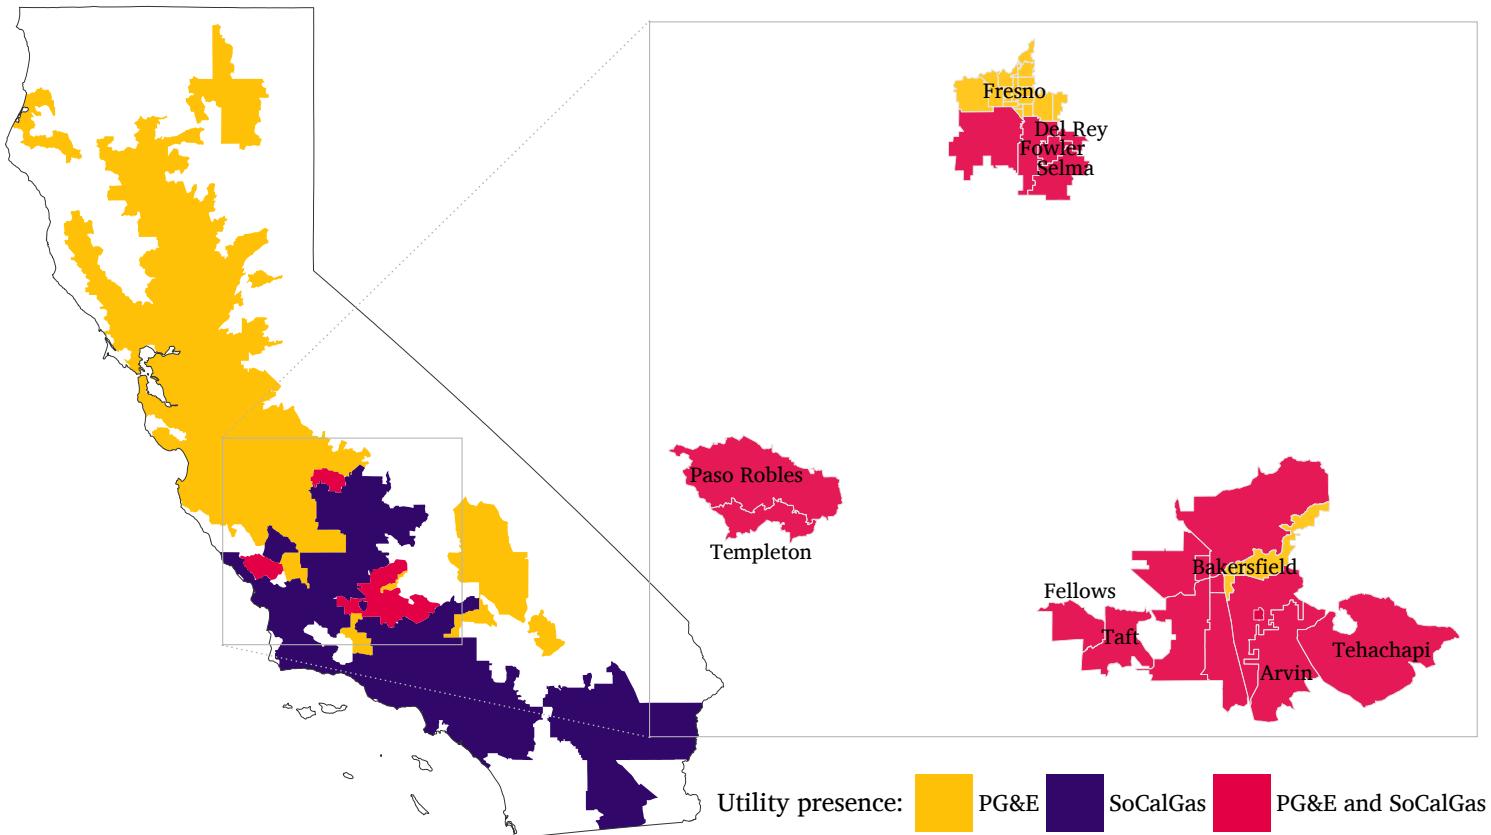
(b) Tier-one daily allowances over time: PG&E (zone R) and SoCalGas (zone 1), 2009–2015



Notes: Households receive daily allowances for baseline (first-tier) consumption as a function of location and season (e.g., climate zone R, January 2009). The household pays the second-tier price on all units that exceed its allowance—comparing total consumption (in the billing period) to total allowance (daily allowance summed across the bills' days).

Figure 5: Natural gas service areas and the study-area discontinuity

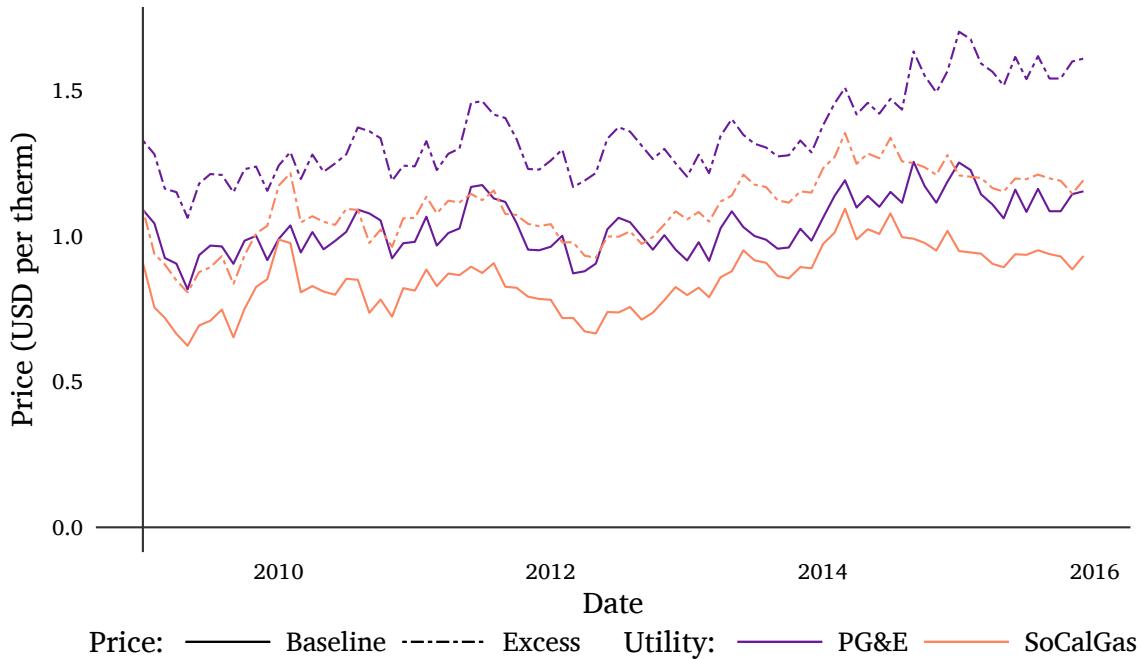
32



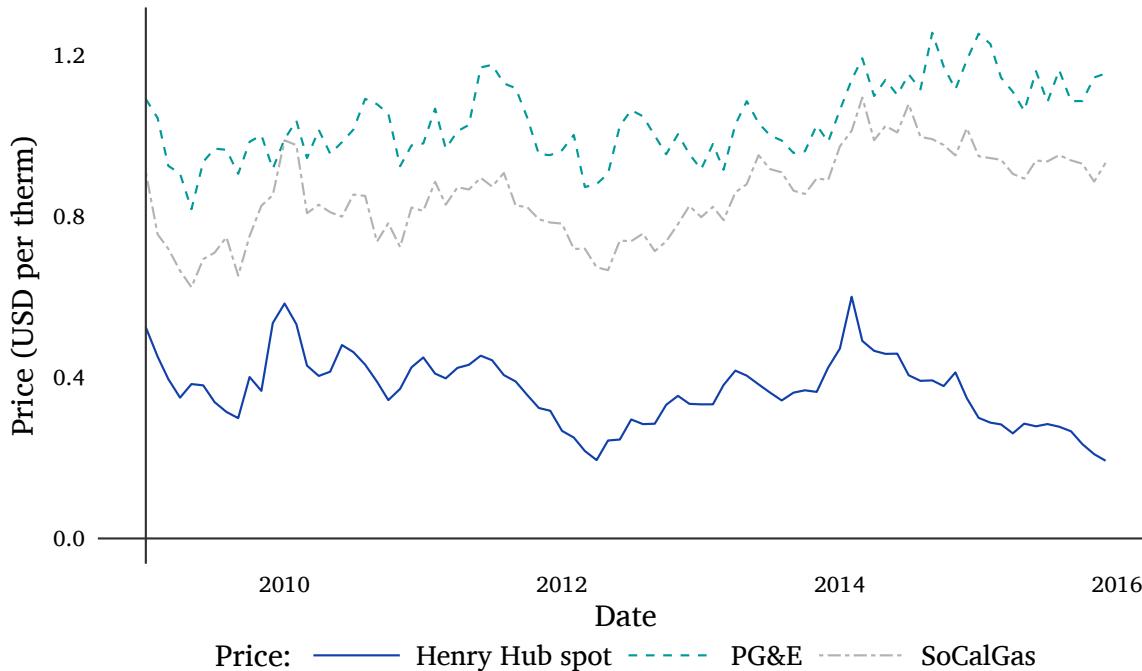
Notes: The left side of the figure displays PG&E's and SoCalGas's services areas (by 5-digit zip code). The right side of the figure zooms in on three clusters of cities that receive service from both utilities. These three clusters of cities encompass 39 zip codes; 18 of these (5-digit) zip codes receive service from both PG&E and SoCalGas. These 18 zip codes represent the main study area for the paper.

Figure 6: Prices across utilities, tiers, and in the spot market, 2009–2015

(a) Price regimes over time: PG&E and SoCalGas, 2009–2015

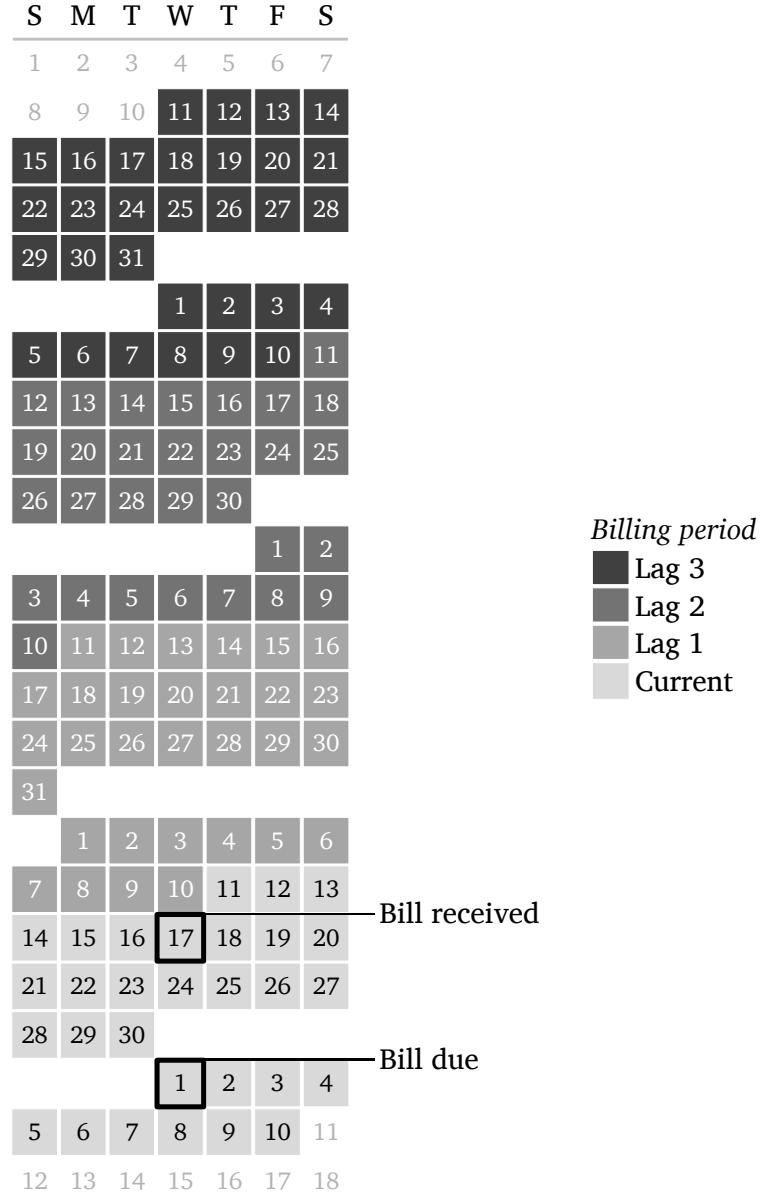


(b) Correlation across prices Three relevant natural gas price series, 2009–2015



Notes: *Baseline* refers to first-tier price, *i.e.*, the price a household pays for its first therm of natural gas. *Excess* refers to the second-tier price, *i.e.*, the price a household pays for each therm that exceeds its first-tier allowance (see Figure 4). The Henry Hub spot price is generally recognized as a national benchmark [U.S. Energy Information Administration 2016a; Levine, Carpenter, and Thapa 2014].

Figure 7: **Calendar months and billing periods:** Four 30-day bills and five months



Notes: The household receives its bill from the *Lag 1* period on the fifth business day of its *current billing period* (the 17th). Payment for the *Lag 1* bill is due two weeks later (on the 1st). Now consider the question “Which lag of price is relevant?”. **Current:** For the household to know the price structure for its current billing period, it must pay attention to the approval status of its utility’s advice-letters correspondence with the CPUC. **Lag 1:** Again, unless the household pays attention to the utility’s CPUC-approved advice letters, it will not know the prices in the *Lag 1* billing period until it receives and opens the bill. The bill arrives several days into the new period, and if the household does not see the bill until payment, it may not learn about the prices of the *Lag 1* bill until the current billing period is nearly complete. Autopay may extend this moment of salience even further into the future. **Lag 2:** Throughout the entirety of the *Current* billing period, the household will know the prices from its *Lag 2* bill, and for a non-zero amount of time, the *Lag 2* bill is likely to be the most recent set of prices the household knows. **Lag 3:** Same level of knowledge as *Lag 2* but less recent.

8 Tables

Table 1: Prior point estimates of the price elasticity of demand for residential natural gas

Paper	Data	Estimate
Davis and Muehlegger (2010)	US state panel	-0.278
Maddala <i>et al.</i> (1997)	US state panel	-0.09 to -0.18
Garcia-Cerrutti (2000)	Calif. county panel	-0.11
Hausman and Kellogg (2015)	US state panel	-0.11
Herbert and Kreil (1989)	Monthly time series	-0.36
Houthakker and Taylor (1970)	Time series	-0.15
Metcalf and Hassett (1999)	RECS HH panel	-0.08 to -0.71
Meier and Rehdanz (2010)	UK HH panel	-0.34 to -0.56
Rehdanz (2007)	Germany HH panel	-0.44 to -0.63

Sources: Authors and Alberini et al. (2011)

Table 2: Billing data summaries

	Full dataset		Border-area dataset	
	PG&E	SoCalGas	PG&E	SoCalGas
N. 5-digit zip codes	597	611	18	18
N. 9-digit zip codes	680,846	610,207	18,047	16,295
N. unique households	5,888,276	2,526,503	152,418	68,407
N. bills	180,663,705	95,335,393	3,401,947	2,352,141
Approx. value (USD)	\$5.71B	\$3.28B	\$120M	\$70.5M

Notes: *Full dataset* refers to all of the PG&E and SoCalGas bills in the data. *Border-area (discontinuity) dataset* refers to the subset of the *full dataset* for households located in the 18 5-digit zip codes served by both utilities during 2010–2014.

Table 3: Summaries of prices, quantities, and other variables of interest

Variable	5% Sample of California			Border-discontinuity sample				
	Split by utility			Overall	Split by season		Split by CARE	
	Overall	PG&E	SoCalGas		Winter	Summer	CARE	Non-Care
Baseline price	0.8901 [0.1686]	0.9823 [0.1206]	0.7432 [0.1242]	0.9026 [0.1419]	0.8836 [0.1361]	0.9204 [0.1448]	0.8080 [0.0854]	0.9811 [0.1311]
Average price	1.0138 [0.1845]	1.1053 [0.1439]	0.8680 [0.1439]	1.0211 [0.1621]	1.0008 [0.1583]	1.0402 [0.1633]	0.9086 [0.1004]	1.1147 [0.1430]
Marginal price	1.0206 [0.2260]	1.1277 [0.186]	0.8500 [0.173]	1.0387 [0.1983]	1.0121 [0.1905]	1.0637 [0.2021]	0.9338 [0.1448]	1.1259 [0.1944]
3 Therms	35.4626 [33.7995]	37.7541 [36.0107]	31.8135 [29.5791]	33.8273 [30.7697]	50.9544 [35.2487]	17.7311 [11.5803]	33.1136 [28.7629]	34.4204 [32.3306]
Days	30.3992 [1.4275]	30.4282 [1.2667]	30.3530 [1.6505]	30.3994 [1.3038]	30.5876 [1.3843]	30.2225 [1.1966]	30.4040 [1.2761]	30.3955 [1.3263]
Therms per day	1.1592 [1.0921]	1.2355 [1.1698]	1.0378 [0.9426]	1.1063 [0.9936]	1.6588 [1.1354]	0.5871 [0.3838]	1.0840 [0.9304]	1.1249 [1.0429]
Total bill	36.8703 [39.5758]	42.3938 [44.0564]	28.0747 [29.0445]	34.9508 [33.8812]	52.0750 [39.8973]	18.8573 [14.0069]	30.3135 [27.2567]	38.8040 [38.1017]
(Percent) CARE	27.43%	26.35%	29.15%	45.38%	45.00%	45.74%	100%	0%

Notes: Unbracketed values provide the variables' means; bracketed values denote the variables' standard deviations. The *5% sample of California* is based upon 5% of PG&E's and SoCalGas's natural gas bills from 2010–2014, sampling at the 5-digit zip code. The *border-discontinuity sample* represents all bills from PG&E and SoCalGas for the 18 5-digit zip codes served by both utilities from 2010–2014.

Table 4: **OLS Results:** Estimating elasticities, varying the dataset, price, and fixed effects

	Dependent variable: Log(Consumption, daily avg.)					
	(1)	(2)	(3)	(4)	(5)	(6)
Log(Marginal price)	0.4698*** (0.0106)	0.4346*** (0.0136)	0.4276*** (0.0134)			
Log(Baseline price)				0.0217 (0.0147)	-0.0918*** (0.0201)	-0.1009*** (0.0209)
Bill HDDs	T	T	T	T	T	T
Household FE	T	T	T	T	T	T
Month-of-sample FE	T	T	F	T	T	F
City by month-of-sample FE	F	F	T	F	F	T
Sample	5% CA	Border	Border	5% CA	Border	Border
N	12,855,910	5,754,088	5,754,088	12,855,910	5,754,088	5,754,088

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Base* or *baseline* price refers to the price the household pays for its first unit (*therm*) of natural gas. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Significance levels:* *10%, **5%, ***1%.

Table 5: First- and second-stage results:
Instrumenting consumers' prices with the Henry Hub spot price

Dependent variable: Log(Consumption, daily avg.)					
	(1) Marginal	(2) Average	(3) Avg. Mrg.	(4) Baseline	(5) Sim. Mrg.
Spot price	0.3679*** (0.0774)	0.3697*** (0.0521)	0.3384*** (0.0570)	0.4699*** (0.0434)	0.3949*** (0.0840)
Spot price × SoCalGas	0.7868*** (0.0299)	0.7174*** (0.0186)	0.9389*** (0.0198)	0.8212*** (0.0176)	0.8174*** (0.0317)
Panel B: Second-stage results					
Log(Price) (instrumented)	-0.2098*** (0.0706)	-0.2312*** (0.076)	-0.1734*** (0.0585)	-0.2030*** (0.065)	-0.1705** (0.0698)
First-stage F stat.	418.4	899.4	1,311.0	1,333.2	369.9
Bill HDDs	T	T	T	T	T
Household FE	T	T	T	T	T
City mo.-of-sample FE	T	T	T	T	T
N	5,754,085	5,754,085	5,754,085	5,754,085	4,682,526

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. (*HH*) *Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Avg.* or *average* price is the total bill divided by quantity. *Avg. Mrg.* or *average marginal* price denotes the quantity-weighted average of the household's marginal price. *Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14). As discussed in the Empirical strategy section, the numbers of observations differ due to the lags required to calculate the *simulated instrument* for marginal price. *Significance levels:* *10%, **5%, ***1%.

Table 6: Second-stage results, comparing lags: Marginal and average prices with HH spot price IV

	Dependent variable: Log(Consumption, daily avg.)					
	Marginal Price			Average Price		
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) No lag	(5) 1 Lag	(6) 2 Lags
Log(Price) <i>instrumented</i>	-0.1121 (0.0762)	-0.0223 (0.0668)	-0.2098*** (0.0706)	-0.1244 (0.0805)	-0.0177 (0.0730)	-0.2312*** (0.0760)
First-stage F stat.	337.9	410.8	418.4	725.8	884.4	899.4
Bill HDDs	T	T	T	T	T	T
Household FE	T	T	T	T	T	T
City month-of-sample FE	T	T	T	T	T	T
N	5,754,088	5,754,088	5,754,085	5,754,088	5,754,088	5,754,085

3

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; etc. Avg. or average price is the total bill divided by quantity. (HH) Spot price refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Significance levels: *10%, **5%, ***1%.

Table 7: Heterogeneity by season or income:
Second-stage results, instrumenting marginal price with HH spot price

Dependent variable: Log(Consumption, daily avg.)

	Marginal Price			
	Split by Season		Split by CARE (Income)	
	(1) Summer	(2) Winter	(3) CARE	(4) Non-CARE
Log(Price) <i>instrumented</i>	0.0519* (0.0285)	-0.3769*** (0.1399)	-0.2443*** (0.0794)	-0.1413** (0.0684)
First-stage F stat.	319.6	174.2	393.7	335.8
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	3,065,917	2,688,168	2,435,135	3,318,950

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Summer* includes April through September. *Winter* includes October through March. CARE households participate in the California Alternative Rates for Energy (CARE) program. CARE targets low-income households and provides a 20-percent discount on natural gas bills. We estimate the heterogeneity results by splitting the sample along the dimension(s) of heterogeneity and then individually estimating the models. *Significance levels:* *10%, **5%, ***1%.

Table 8: Heterogeneity by season and income:
Second-stage results, instrumenting marginal price with HH spot price

Dependent variable: Log(Consumption, daily avg.)

	Marginal Price			
	(1) Summer CARE	(2) Summer Non-CARE	(3) Winter CARE	(4) Winter Non-CARE
Log(Price) <i>instrumented</i>	0.0457 (0.0353)	0.0742** (0.0324)	-0.5226*** (0.1424)	-0.3173** (0.1498)
First-stage F stat.	303.4	237.1	145.6	156.7
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	1,293,144	1,772,773	1,141,991	1,546,177

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Summer* includes April through September. *Winter* includes October through March. CARE households participate in the California Alternative Rates for Energy (CARE) program. CARE targets low-income households and provides a 20-percent discount on natural gas bills. We estimate the heterogeneity results by splitting the sample along the dimension(s) of heterogeneity and then individually estimating the models. *Significance levels:* *10%, **5%, ***1%.

Table of contents

1	Introduction	1
2	Institutional setting	4
2.1	Market segments	4
2.2	Natural gas pricing in California	6
3	Data	8
3.1	Natural gas billing data	8
3.2	Weather data	9
4	Empirical strategy	9
4.1	Estimating equation	10
4.2	Challenges	10
4.3	Solutions	12
4.3.1	Discontinuities	12
4.3.2	Instrumental variables	14
4.3.3	Instrumented prices and simulated instruments	16
5	Results	19
5.1	Pooled price elasticity of demand for natural gas	19
5.2	Heterogeneity	21
5.2.1	Income heterogeneity	21
5.2.2	Seasonal heterogeneity	22
5.2.3	Income-by-season heterogeneity	22
6	Conclusion	23
7	Figures	28
8	Tables	35
Descriptive tables	35	
2SLS results	38	
Heterogeneity results	40	
Table of contents	42	
A	Appendix	45
A.1	Appendix figures	45
A.2	Appendix tables	48
A.3	Calculating bills	57
A.4	Data work	58

List of Figures

1	U.S. Natural gas institutional organization	28
2	U.S. natural gas pipeline network	29
3	Henry Hub natural gas spot price: Daily, 1997–2016	30
4	Households' allowances and prices	31
5	Natural gas service areas and the study-area discontinuity	32
6	Prices across utilities, tiers, and in the spot market, 2009–2015	33
7	Calendar months and billing periods: Four 30-day bills and five months	34
A1	California's 16 CEC climate zones determine daily allowance within season	45
A2	Example bill: PG&E residential natural gas bill	46
A3	PRISM: Mean temperature raster for 15 June 2010	47

List of Tables

1	Prior point estimates of the price elasticity of demand for residential natural gas	35
2	Billing data summaries	35
3	Summaries of prices, quantities, and other variables of interest	36
4	OLS Results: Estimating elasticities, varying the dataset, price, and fixed effects	37
5	First- and second-stage results: Instrumenting consumers' prices with the Henry Hub spot price	38
6	Second-stage results, comparing lags: Marginal and average prices with HH spot price IV	39
7	Heterogeneity by season or income: Second-stage results, instrumenting marginal price with HH spot price	40
8	Heterogeneity by season and income: Second-stage results, instrumenting marginal price with HH spot price	41
A1	Price correlation: Bivariate correlations between types of prices	48
A2	Testing the simulated instrument: Regressing marginal price on <i>simulated</i> marginal price	48
A3	Second-stage results, comparing lags: Marginal price with HH spot price IV	49
A4	Second-stage results, comparing lags: Sim. marginal price with HH spot price IV	49
A5	Second-stage results, comparing lags: Avg. marginal price with HH spot price IV	50
A6	Second-stage results, comparing lags: Average price with HH spot price IV	50
A7	Second-stage results, comparing lags: Baseline price with HH spot price IV	51
A8	Heterogeneity by season or income: Second-stage results, instrumenting average price with HH spot price	52

A9 Heterogeneity by season and income:	
Second-stage results, instrumenting average price with HH spot price	53
A10 First-stage results:	
Robustness to specification: Marginal price instrumented with spot price	54
A11 Second-stage results:	
Robustness to specification: Marginal price instrumented with spot price	54
A12 Second-stage results:	
Robustness to specification: Average price instrumented with spot price	55
A13 Second-stage results:	
Robustness to specification: Avg. mrg. price instrumented with spot price	55
A14 Second-stage results:	
Robustness to specification: Baseline price instrumented with spot price	56
A15 Second-stage results:	
Robustness to specification: Baseline price instrumented with spot price	56
A16 Billing data description: Columns within the billing data	57

A Appendix

A.1 Appendix figures

Figure A1: California's 16 CEC climate zones determine daily allowance within season



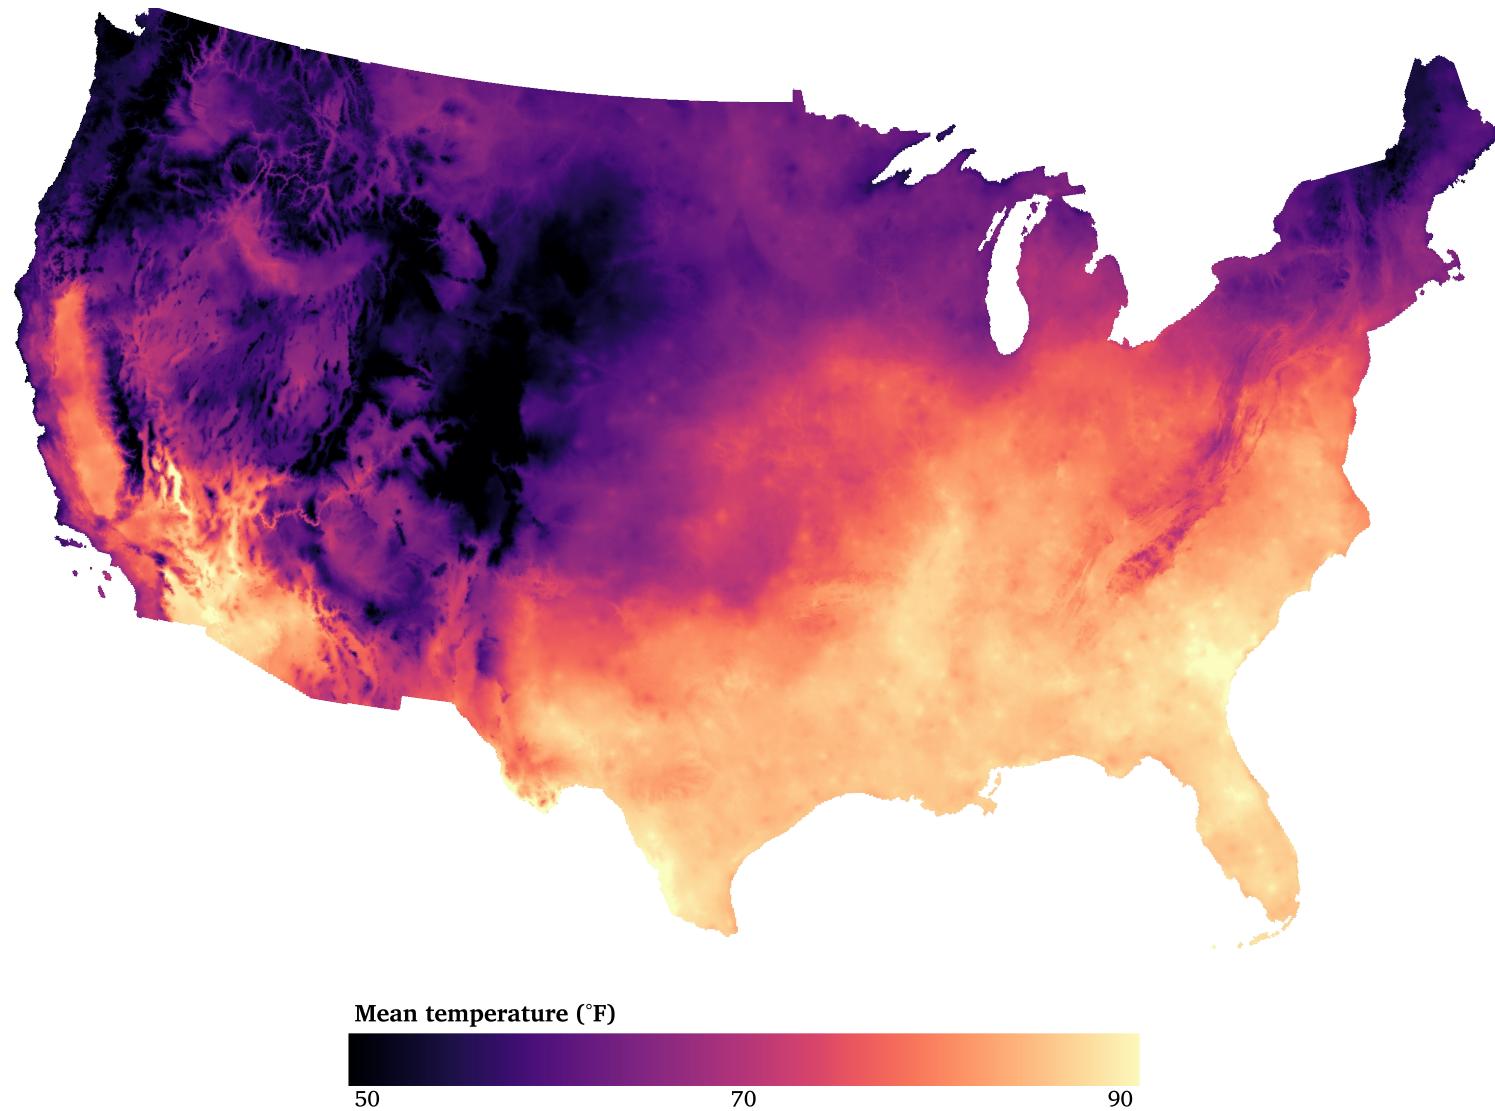
Notes: The shapefile underlying this map comes from the [California Energy Commission \(CEC\)](#). This map constitutes the CEC's climate-based building zones, which affect a number of energy policies, including households' baseline allowances. California Energy Commission [2015, 2017](#)

Figure A2: Example bill: PG&E residential natural gas bill

 ENERGY STATEMENT www.pge.com/MyEnergy	Account No: [REDACTED] Statement Date: 12/25/2016 Due Date: 01/17/2017								
Details of Gas Charges 11/24/2016 - 12/23/2016 (30 billing days)									
Service For: [REDACTED] Service Agreement ID: [REDACTED] Rate Schedule: G1 T Residential Service									
11/24/2016 – 11/30/2016 Your Tier Usage <table border="1" style="display: inline-table;"><tr><td style="width: 20px; height: 15px;"></td><td style="width: 20px; height: 15px; background-color: #ccc;"></td></tr></table>									
▼ <table border="0" style="width: 100%;"> <tr> <td style="width: 30%;">Tier 1 Allowance</td> <td style="width: 40%;">12.53 Therms (7 days x 1.79 Therms/day)</td> </tr> <tr> <td>Tier 1 Usage</td> <td>2.33330 Therms @ \$1.42692 \$3.33</td> </tr> <tr> <td>Gas PPP Surcharge (\$0.10197 /Therm)</td> <td>0.24</td> </tr> <tr> <td>Berkeley Utility Users' Tax (7.500%)</td> <td>0.25</td> </tr> </table>		Tier 1 Allowance	12.53 Therms (7 days x 1.79 Therms/day)	Tier 1 Usage	2.33330 Therms @ \$1.42692 \$3.33	Gas PPP Surcharge (\$0.10197 /Therm)	0.24	Berkeley Utility Users' Tax (7.500%)	0.25
Tier 1 Allowance	12.53 Therms (7 days x 1.79 Therms/day)								
Tier 1 Usage	2.33330 Therms @ \$1.42692 \$3.33								
Gas PPP Surcharge (\$0.10197 /Therm)	0.24								
Berkeley Utility Users' Tax (7.500%)	0.25								
12/01/2016 – 12/23/2016 Your Tier Usage <table border="1" style="display: inline-table;"><tr><td style="width: 20px; height: 15px;"></td><td style="width: 20px; height: 15px; background-color: #ccc;"></td></tr></table>									
▼ <table border="0" style="width: 100%;"> <tr> <td style="width: 30%;">Tier 1 Allowance</td> <td style="width: 40%;">41.17 Therms (23 days x 1.79 Therms/day)</td> </tr> <tr> <td>Tier 1 Usage</td> <td>7.666670 Therms @ \$1.36245 \$10.45</td> </tr> <tr> <td>Gas PPP Surcharge (\$0.10197 /Therm)</td> <td>0.78</td> </tr> <tr> <td>Berkeley Utility Users' Tax (7.500%)</td> <td>0.78</td> </tr> </table>		Tier 1 Allowance	41.17 Therms (23 days x 1.79 Therms/day)	Tier 1 Usage	7.666670 Therms @ \$1.36245 \$10.45	Gas PPP Surcharge (\$0.10197 /Therm)	0.78	Berkeley Utility Users' Tax (7.500%)	0.78
Tier 1 Allowance	41.17 Therms (23 days x 1.79 Therms/day)								
Tier 1 Usage	7.666670 Therms @ \$1.36245 \$10.45								
Gas PPP Surcharge (\$0.10197 /Therm)	0.78								
Berkeley Utility Users' Tax (7.500%)	0.78								
Total Gas Charges \$15.83									
Service Information									
Meter # [REDACTED] Current Meter Reading 3,036 Prior Meter Reading 3,026 Difference 10 Multiplier 1.046971 Total Usage 10.000000 Therms Baseline Territory T Serial C									
Gas Procurement Costs (\$/Therm)									
11/24/2016 - 11/30/2016 \$0.45875 12/01/2016 - 12/23/2016 \$0.39428									

Notes: This 30-day bill for a PG&E customer (one of the authors) overlaps two calendar months in 2016: 7 days in November (24–30) and 23 days December (01–23). Because PG&E's prices vary with the calendar month, PG&E needs to split consumption by calendar month. To achieve this task, PG&E assumes the customer consumed evenly across all days in the bill. Specifically, PG&E calculates that the customer consumed 10 therms and assigns the same amount of consumption to each day during the 30-day period. Thus, PG&E assigns $10 \times 7/30 \approx 2.33$ to November (the consumer spent 7 days in November in this 30-day bill) and $10 \times 23/30 \approx 7.67$ to December (the consumer spent 23 days in November in this 30-day bill).

Figure A3: PRISM: Mean temperature raster for 15 June 2010



A.2 Appendix tables

Table A1: **Price correlation:** Bivariate correlations between types of prices

		Type of Price			
	Marginal	Average	Avg. Mrg.	Baseline	Sim. mrg.
Marginal	1				
Average	0.8898	1			
Avg. Mrg.	0.8628	0.9421	1		
Baseline	0.7901	0.942	0.9202	1	
Sim. mrg.	0.8503	0.849	0.8174	0.781	1

Notes: *Avg.* or *average* price is the total bill divided by quantity. *Avg. Mrg.* or *average marginal* price denotes the quantity-weighted average of the household's marginal price. *Base* or *baseline* price refers to the price the household pays for its first unit (*therm*) of natural gas. *Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14).

Table A2: **Testing the simulated instrument:**
Regressing marginal price on *simulated* marginal price

Dependent variable: Marginal price		
	(1)	(2)
Simulated marginal price	0.6425*** (0.00435)	0.6444*** (0.00433)
Bill HDDs	T	T
Household FE	T	T
City month-of-sample FE	T	T
Lags used for sim. inst.	10–14	11–13
N	4,892,064	4,785,877

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14 or 11 through 13). As discussed in the Empirical strategy section, the numbers of observations differ due to the lags required to calculate the *simulated instrument* for marginal price. *Significance levels:* *10%, **5%, ***1%.

Table A3: Second-stage results, comparing lags: Marginal price with HH spot price IV

Dependent variable: Log(Consumption, daily avg.)

	Lag of Marginal Price			
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) 3 Lags
Log(Price) <i>instrumented</i>	-0.1121 (0.0762)	-0.0223 (0.0668)	-0.2098*** (0.0706)	-0.1582** (0.0698)
First-stage F stat.	337.9	410.8	418.4	403.4
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	5,754,088	5,754,088	5,754,085	5,754,079

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; etc. (*HH*) *Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill.
Significance levels: *10%, **5%, ***1%.

Table A4: Second-stage results, comparing lags: Sim. marginal price with HH spot price IV

Dependent variable: Log(Consumption, daily avg.)

	Lag of Simulated Marginal Price			
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) 3 Lags
Log(Price) <i>instrumented</i>	-0.0549 (0.0718)	0.0329 (0.0626)	-0.1705** (0.0698)	-0.1596** (0.0720)
First-stage F stat.	379.6	393.2	369.9	332.1
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	4,892,064	4,785,877	4,682,526	4,590,790

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; etc. *Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14). (*HH*) *Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Significance levels:* *10%, **5%, ***1%.

Table A5: **Second-stage results, comparing lags:** Avg. marginal price with HH spot price IV

Dependent variable: Log(Consumption, daily avg.)

	Lag of Average Marginal Price			
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) 3 Lags
Log(Price) <i>instrumented</i>	-0.0853 (0.0618)	-0.0313 (0.0568)	-0.1734*** (0.0585)	-0.1356** (0.0585)
First-stage F stat.	1,036.4	1,275.3	1,311.0	1,306.1
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	5,754,088	5,754,088	5,754,085	5,754,079

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; *etc. Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14). *(HH) Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Significance levels:* *10%, **5%, ***1%.

Table A6: **Second-stage results, comparing lags:** Avgerage price with HH spot price IV

Dependent variable: Log(Consumption, daily avg.)

	Lag of Average Price			
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) 3 Lags
Log(Price) <i>instrumented</i>	-0.1244 (0.0805)	-0.0177 (0.0730)	-0.2312*** (0.0760)	-0.1680** (0.0749)
First-stage F stat.	725.8	884.4	899.4	923.7
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	5,754,088	5,754,088	5,754,085	5,754,079

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; *etc. Avg.* or *average* price is the total bill divided by quantity. *(HH) Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Significance levels:* *10%, **5%, ***1%.

Table A7: **Second-stage results, comparing lags:** Baseline price with HH spot price IV

Dependent variable: Log(Consumption, daily avg.)

	Lag of Baseline Price			
	(1) No lag	(2) 1 Lag	(3) 2 Lags	(4) 3 Lags
Log(Price) <i>instrumented</i>	-0.1164* (0.0684)	-0.0066 (0.0637)	-0.2030*** (0.0650)	-0.1396** (0.0630)
First-stage F stat.	1,143.4	1,241.8	1,333.2	1,533.2
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	5,754,088	5,754,088	5,754,085	5,754,079

Notes: With regards to lags: *No lag* refers to the price for the household's contemporaneous bill; *1 Lag* refers to the price on the household's previous bill; etc. *Base* or *baseline* price refers to the price the household pays for its first unit (*therm*) of natural gas. (*HH*) *Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Significance levels:* *10%, **5%, ***1%.

Table A8: **Heterogeneity by season or income:**
 Second-stage results, instrumenting average price with HH spot price

Dependent variable: Log(Consumption, daily avg.)

	Average Price			
	Split by Season		Split by CARE (Income)	
	(1) Summer	(2) Winter	(3) CARE	(4) Non-CARE
Log(Price) <i>instrumented</i>	0.0579* (0.0316)	-0.4694*** (0.1586)	-0.2650*** (0.0834)	-0.1557** (0.0740)
First-stage F stat.	765.7	223.4	814.7	745.8
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	3,065,917	2,688,168	2,435,135	3,318,950

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. *Summer* includes April through September. *Winter* includes October through March. CARE households participate in the California Alternative Rates for Energy (CARE) program. CARE targets low-income households and provides a 20-percent discount on natural gas bills. We estimate the heterogeneity results by splitting the sample along the dimension(s) of heterogeneity and then individually estimating the models. Avg. or average price is the total bill divided by quantity. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. Significance levels: *10%, **5%, ***1%.

Table A9: Heterogeneity by season and income:
 Second-stage results, instrumenting average price with HH spot price

Dependent variable: Log(Consumption, daily avg.)

	Average Price			
	(1) Summer CARE	(2) Summer Non-CARE	(3) Winter CARE	(4) Winter Non-CARE
Log(Price) <i>instrumented</i>	0.0495 (0.0384)	0.0828** (0.0359)	-0.6106*** (0.1570)	-0.3971** (0.1687)
First-stage F stat.	691.5	591.9	212.7	184.8
Bill HDDs	T	T	T	T
Household FE	T	T	T	T
City month-of-sample FE	T	T	T	T
N	1,293,144	1,772,773	1,141,991	1,546,177

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Summer* includes April through September. *Winter* includes October through March. CARE households participate in the California Alternative Rates for Energy (CARE) program. CARE targets low-income households and provides a 20-percent discount on natural gas bills. We estimate the heterogeneity results by splitting the sample along the dimension(s) of heterogeneity and then individually estimating the models. Avg. or average price is the total bill divided by quantity. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. Significance levels: *10%, **5%, ***1%.

Table A10: **First-stage results:**

Robustness to specification: Marginal price instrumented with spot price

Dependent variable: Log(Marginal price)				
	(1)	(2)	(3)	(4)
Spot price	0.3398*** (0.0757)	0.3679*** (0.0774)	0.3806*** (0.0798)	0.3955*** (0.0547)
Spot price × SoCalGas	0.7858*** (0.0300)	0.7868*** (0.0299)	0.7856*** (0.0302)	0.7385*** (0.0378)
First-stage F stat.	416.1	418.4	415.2	367.0
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	5,754,085	5,754,085	5,754,085	5,754,085

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. (*HH*) *Spot price* refers to the weekly average spot price for natural gas at Louisiana's Henry Hub in the week preceding the utility's price change. *Significance levels:* *10%, **5%, ***1%.

Table A11: **Second-stage results:**

Robustness to specification: Marginal price instrumented with spot price

Dependent variable: Log(Consumption, daily avg.)				
	(1)	(2)	(3)	(4)
Log(Marginal price) instrumented	-0.3623*** (0.0854)	-0.2098*** (0.0706)	-0.1705*** (0.0621)	-0.1495** (0.063)
First-stage F stat.	416.1	418.4	415.2	367.0
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	5,754,085	5,754,085	5,754,085	5,754,085

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Significance levels:* *10%, **5%, ***1%.

Table A12: **Second-stage results:**

Robustness to specification: Average price instrumented with spot price

Dependent variable: Log(Consumption, daily avg.)

	(1)	(2)	(3)	(4)
Log(Average price) <i>instrumented</i>	-0.4076*** (0.0911)	-0.2312*** (0.076)	-0.1891*** (0.067)	-0.1574** (0.0656)
First-stage F stat.	897.5	899.4	881.1	661.1
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	5,754,085	5,754,085	5,754,085	5,754,085

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Avg.* or *average price* is the total bill divided by quantity. *Significance levels:* *10%, **5%, ***1%.

Table A13: **Second-stage results:**

Robustness to specification: Avg. mrg. price instrumented with spot price

Dependent variable: Log(Consumption, daily avg.)

	(1)	(2)	(3)	(4)
Log(Avg. marginal price) <i>instrumented</i>	-0.2951*** (0.0697)	-0.1734*** (0.0585)	-0.1529*** (0.0514)	-0.1330** (0.0549)
First-stage F stat.	1,299.9	1,311.0	1,275.8	780.6
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	5,754,085	5,754,085	5,754,085	5,754,085

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Avg. Mrg.* or *average marginal price* denotes the quantity-weighted average of the household's marginal price. *Significance levels:* *10%, **5%, ***1%.

Table A14: **Second-stage results:**

Robustness to specification: Baseline price instrumented with spot price

Dependent variable: Log(Consumption, daily avg.)

	(1)	(2)	(3)	(4)
Log(Simulated mrg. price) <i>instrumented</i>	-0.3148*** (0.0843)	-0.1705** (0.0698)	-0.1310** (0.0602)	-0.1025 (0.0675)
First-stage F stat.	368.9	369.9	331.3	181.9
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	4,682,526	4,682,526	4,682,526	4,682,526

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Sim. Mrg.* or *simulated marginal* price is the household's marginal price (using the relevant pricing regime) as a function of the household's historical consumption patterns (lagged bills 10 through 14). *Significance levels:* *10%, **5%, ***1%.

Table A15: **Second-stage results:**

Robustness to specification: Baseline price instrumented with spot price

Dependent variable: Log(Consumption, daily avg.)

	(1)	(2)	(3)	(4)
Log(Baseline price) <i>instrumented</i>	-0.3643*** (0.077)	-0.2030*** (0.065)	-0.1653*** (0.0576)	-0.1376** (0.0572)
First-stage F stat.	1,322.9	1,333.2	1,187.3	762.5
Bill HDDs	F	T	T	T
Household FE	T	T	T	T
City by month-of-sample FE	T	T	F	F
City by week-of-sample FE	F	F	T	F
Zip by week-of-sample FE	F	F	F	T
N	5,754,085	5,754,085	5,754,085	5,754,085

Notes: Each column denotes a separate regression. Errors are two-way clustered within (1) household and (2) utility by climate-zone by billing-cycle (the level at which price varies). All regressions include heating degree days (HDDs) within the households' bill. Each price in the table is the second lag of price, *i.e.*, the prices from two bills prior to the current bill. *Base* or *baseline* price refers to the price the household pays for its first unit (*therm*) of natural gas. *Significance levels:* *10%, **5%, ***1%.

Table A16: **Billing data description:** Columns within the billing data

Feature name	Description
Account ID	Unique identifier for household account with the utility
Premise ID	Unique physical-location based identifier
Prior read date	Effectively the start date of the bill
Current read date	Effectively the end date of the bill
Gas rate schedule	Classifies type of customer (and the customer's price regime)
Gas usage	Volume of gas consumed during billing period (in therms)
Bill revenue	Total bill charged to household for the current billing period
Climate band	California Public Utility Commission-based climate region
Service address 9-digit zip	Household's 9-digit zip code
Service start date	Date on which the household began service
Service stop date	Date on which the household ended service

A.3 Calculating bills

As discussed in the body of the paper, the majority of bills do not line up with calendar months. Consequently, households' billing periods do not line up with utilities' monthly changes in price (or with changes in daily allowances—resulting from changes in seasons). Thus single bill will typically span multiple price regimes. The two utilities deal with change in price in subtly different ways. This “problem” results from the fact that neither utility knows households' *daily* consumption.

PG&E When a PG&E customer's bill spans multiple calendar months (price regimes), PG&E calculates individual bills for each month. However, because PG&E does not know the daily consumption levels, they assume constant daily consumption throughout the billing period.

SoCalGas In the case that a SoCalGas customer’s bill spans multiple calendar months (price regimes), SoCalGas computes time-weighted average prices (and allowances)—aggregating the prices and allowances from the calendar months by the number of days the bill spent in each month.

A.4 Data work

In this section, we describe the exclusion and cleaning choices that we made while preparing the data for analysis. Our R scripts are available upon request—though the data themselves cannot be shared due to agreements with the utilities and the IRB.

Exclusions:

- We omitted SoCalGas price data from advice letters 3644, 3680, 3695, 3807, 4053, and 4061, as they were updated by letters 3660, 3697, 3697, 3810, 4055, and 4069, respectively.
- We dropped pre-2008 data (PG&E and prices/allowances), as SoCalGas did not share billing data for pre-2009 bills.
- We trimmed the shortest 2.5% and longest 2.5% bills (resulted in keeping bills of length between 28-34 days). We did this to omit the first or last bills for a household and bills that were irregular for any other reasons.
- We dropped bills missing any critical information: number of therms (quantity), revenue, etc. .
- We dropped **households** whose bills (revenue or quantity) were outside the central 99% of data (*i.e.*, in the bottom 0.5% or in the top 0.5%).
- We dropped bills whose total revenue we could not predict within five percent (using known prices, quantities, and discounts).
- We dropped bills for exactly **zero** therms.

CARE status While the datasets presumably denoted CARE (California Alternate Rates for Energy) households, we found many households not denoted as CARE households whose

charges were consistent with CARE pricing (*i.e.*, charges were 80 percent of the standard tariffs). We classified these households as CARE households.